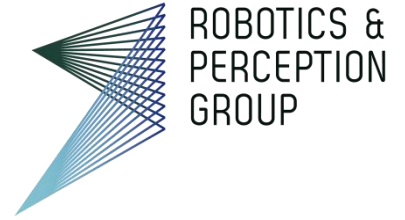




University of  
Zurich<sup>UZH</sup>

**ETH** zürich

Institute of Informatics – Institute of Neuroinformatics



# Lecture 04

## Image Filtering

Davide Scaramuzza

<http://rpg.ifi.uzh.ch>

# No exercise this afternoon

Date	Time	Description of the lecture/exercise	Lecturer
20.09.2018	10:15 - 12:00	01 – Introduction	Davide Scaramuzza
27.09.2018	10:15 - 12:00	02 - Image Formation 1: perspective projection and camera models	Davide Scaramuzza
	13:15 – 15:00	Exercise 1: Augmented reality wireframe cube	Titus Cieslewski & Mathias Gehrig
04.10.2018	10:15 - 12:00	03 - Image Formation 2: camera calibration algorithms	Guillermo Gallego
	13:15 – 15:00	Exercise 2: PnP problem	Antonio Loquercio & Mathias Gehrig
1.10.2018	10:15 - 12:00	04 - Filtering & Edge detection	Davide Scaramuzza
18.10.2018	10:15 - 12:00	05 - Point Feature Detectors 1: Harris detector	Guillermo Gallego
	13:15 – 15:00	Exercise 3: Harris detector + descriptor + matching	Antonio Loquercio & Mathias Gehrig
25.10.2018	10:15 - 12:00	06 - Point Feature Detectors 2: SIFT, BRIEF, BRISK	Davide Scaramuzza
01.11.2018	10:15 - 12:00	07 - Multiple-view geometry 1	Guillermo Gallego
	13:15 – 15:00	Exercise 4: Stereo vision: rectification, epipolar matching, disparity, triangulation	Antonio Loquercio & Mathias Gehrig
08.11.2018	10:15 - 12:00	08 - Multiple-view geometry 2	Davide Scaramuzza
	13:15 – 15:00	Exercise 5: Eight-Point algorithm	Antonio Loquercio & Mathias Gehrig
15.11.2018	10:15 - 12:00	09 - Multiple-view geometry 3	Davide Scaramuzza
	13:15 – 15:00	Exercise 6: P3P algorithm and RANSAC	Antonio Loquercio & Mathias Gehrig
22.11.2018	10:15 - 12:00	10 - Dense 3D Reconstruction (Multi-view Stereo)	Davide Scaramuzza
	13:15 – 15:00	Exercise session: Intermediate VO Integration	Antonio Loquercio & Mathias Gehrig
29.11.2018	10:15 - 12:00	11 - Optical Flow and Tracking (Lucas-Kanade)	Davide Scaramuzza
	13:15 – 15:00	Exercise 7: Lucas-Kanade tracker	Antonio Loquercio & Mathias Gehrig
06.12.2018	10:15 - 12:00	12 – Place recognition	Davide Scaramuzza
	13:15 – 15:00	Exercise session: Deep Learning Tutorial	Antonio Loquercio
13.12.2018	10:15 - 12:00	13 – Visual inertial fusion	Davide Scaramuzza
	13:15 – 15:00	Exercise 8: Bundle adjustment	Antonio Loquercio & Mathias Gehrig
20.12.2018	10:15 - 12:00	14 - Event based vision	Davide Scaramuzza
	12:30 – 13:30	<b>Scaramuzza's lab visit and live demonstrations:</b> Andreasstrasse 15, 2.11, 8050	Davide Scaramuzza & his lab
	14:00 – 16:00	Exercise session: final VO integration (it will take place close to Scaramuzza's lab)	Antonio Loquercio & Mathias Gehrig

# Image filtering

- The word *filter* comes from frequency-domain processing, where “filtering” refers to the process of accepting or rejecting certain frequency components
- We distinguish between low-pass and high-pass filtering
  - A **low-pass filter** smooths an image (retains low-frequency components)
  - A **high-pass filter** retains the contours (also called edges) of an image (high frequency)

Low-pass filtered image



High-pass filtered image



# Low-pass filtering

# Low-pass filtering

## Motivation: noise reduction

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian distribution



Original



Salt and pepper noise




Impulse noise



Gaussian noise

# Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$


How could we reduce the noise to try to recover the "ideal image"?

# Moving average

- Replaces each pixel with an average of all the values in its neighborhood
- Assumptions:
  - Expect pixels to be like their neighbors
  - Expect noise process to be independent from pixel to pixel

# Moving average


- Replaces each pixel with an average of all the values in its neighborhood
- Moving average in 1D:






# Weighted Moving Average

- Can add weights to our moving average
- *Weights*  $[1, 1, 1, 1, 1] / 5$



# Weighted Moving Average


- Non-uniform weights  $[1, 4, 6, 4, 1] / 16$



# This operation is called *convolution*

Example of convolution of two sequences (or “signals”)

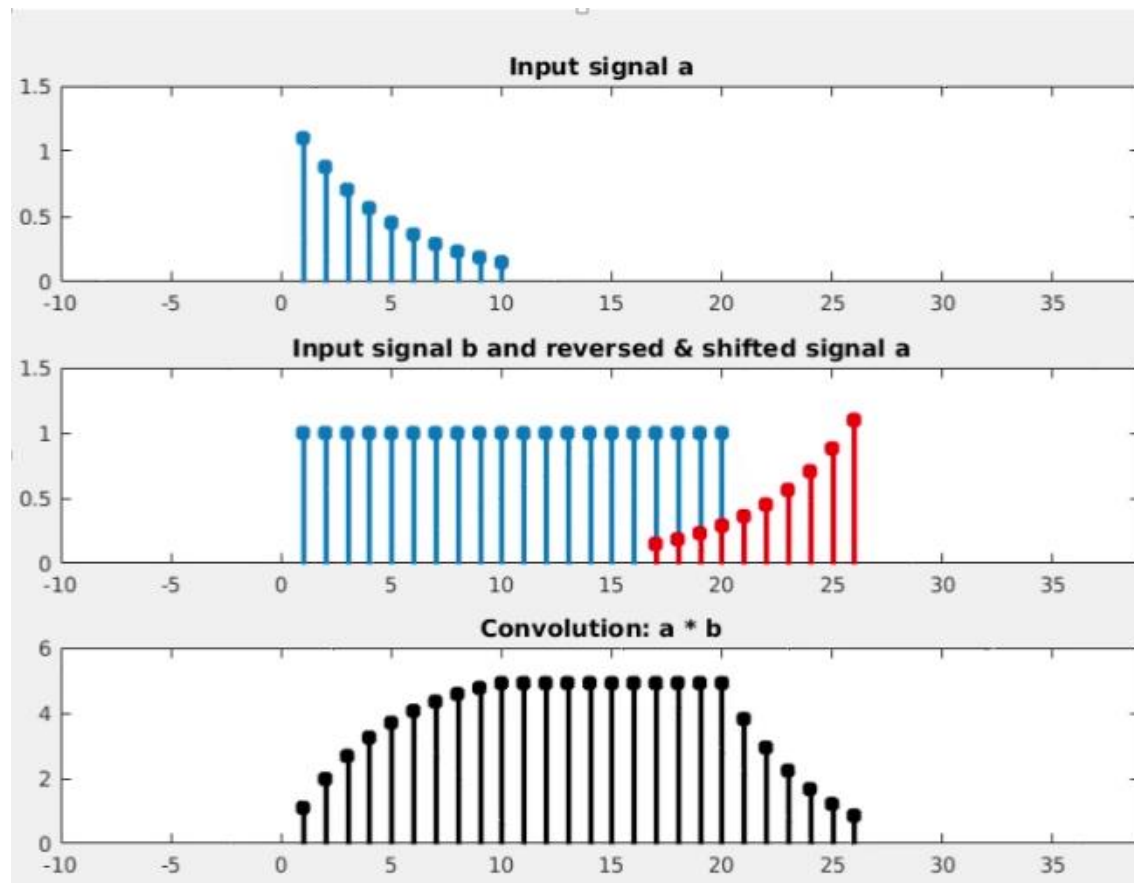
- One of the sequences is flipped (right to left) before sliding over the other
- Notation:  $a \star b$
- Nice properties: linearity, associativity, commutativity, etc.



# This operation is called *convolution*

Example of convolution of two sequences (or “signals”)

- One of the sequences is flipped (right to left) before sliding over the other
- Notation:  $a * b$
- Nice properties: linearity, associativity, commutativity, etc.




# 2D Filtering

- Convolution:
  - Flip the filter in both dimensions (bottom to top, right to left) (=180 deg turn)
  - Then slide the filter over the image and compute sum of products

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k F[x - u, y - v] H[u, v]$$

$$G = F * H$$



Filtering an image: replace each pixel with a linear combination of its neighbors.

The **filter**  $H$  is also called “**kernel**” or “**mask**”.

# Review: Convolution vs. Cross-correlation

## Convolution

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k F[x - u, y - v] H[u, v]$$

$$G = F * H$$

Properties: linearity, associativity, commutativity

## Cross-correlation

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k F[x + u, y + v] H[u, v]$$

$$G = F \otimes H$$

Properties: linearity, but not associativity and commutativity

For a Gaussian or box filter, how will the outputs differ?

# Example: Moving Average in 2D

Input image

Filtered image

$F[x, y]$

$G[x, y]$

“box filter”

$\frac{1}{9}$

1	1	1	0	0	0	0	0	0	0
1	1	1	0	0	0	0	0	0	0
1	1	1	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0								

# Example: Moving Average in 2D

Input image

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$G[x, y]$

	0	10							



# Example: Moving Average in 2D

Input image

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$G[x, y]$

	0	10	20						

# Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

	0	10	20	30					

# Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

	0	10	20	30	30				

# Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	


# Example: Moving Average in 2D



Box filter:  
white = high value, black = low value



original



filtered

# Gaussian filter

- What if we want the closest pixels to have higher influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$F[x, y]$

1	2	1
2	4	2
1	2	1



$H[u, v]$

This kernel is the approximation of a Gaussian function:



$$H[u, v] = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



# Smoothing with a Gaussian




# Compare the result with a box filter






# Gaussian filters

- What parameters matter?
- **Size** of the kernel
  - NB: a Gaussian function has **infinite support**, but discrete filters use finite kernels



$\sigma = 5$  pixels  
with 10 x 10 pixel kernel




$\sigma = 5$  pixels  
with 30 x 30 pixel kernel


Which one is better (left or right)?

# Gaussian filters

- What parameters matter?
- **Variance** of Gaussian: control the amount of smoothing



$\sigma = 2$  pixels  
with 30 x 30 pixel kernel





$\sigma = 5$  pixels  
with 30 x 30 pixel kernel


Recall: standard deviation =  $\sigma$  [pixels], variance =  $\sigma^2$  [pixels<sup>2</sup>]

# Smoothing with a Gaussian

$\sigma$  is called “scale” or “width” or “spread” of the Gaussian kernel, and controls the amount of smoothing.




...




# Sample Matlab code

```
>> hsize = 20;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```


```
>> mesh(h);
```



```
>> imagesc(h);
```



```
>> im = imread('panda.jpg');  
>> outim = imfilter(im, h);  
>> imshow(outim);
```



outim 28

# Boundary issues

- What about near the image edges?
  - the filter window falls off the edges of the image
  - need to pad the image borders
  - methods:
    - zero padding (black)
    - wrap around
    - copy edge
    - reflect across edge



# Summary on (linear) smoothing filters

- Smoothing filter
  - has positive values (also called coefficients)
  - sums to 1  $\rightarrow$  preserve brightness of constant regions
  - removes “high-frequency” components; “low-pass” filter

# Non-linear filtering

# Effect of smoothing filters

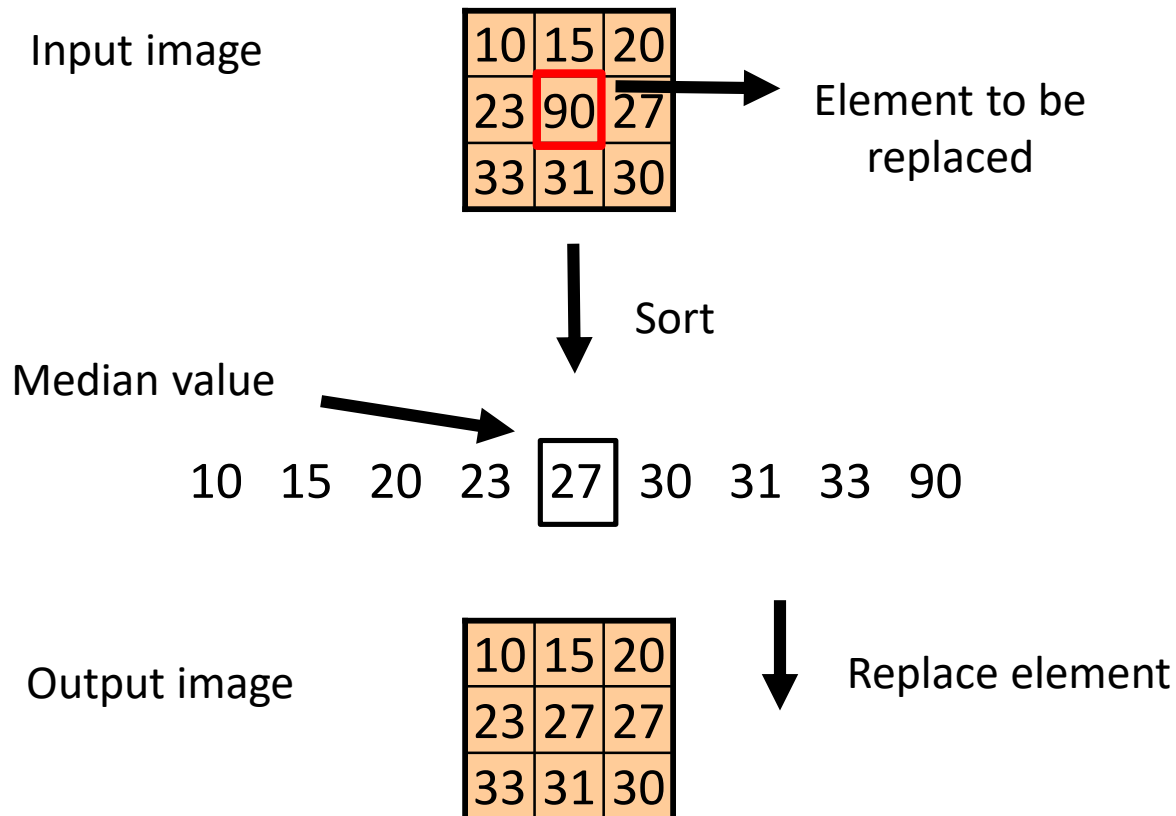


Linear smoothing filters do not alleviate salt and pepper noise!




# Median filter

- It is a non-linear filter
- Removes spikes: good for “*impulse noise*” and “*salt & pepper noise*”





# Median filter

Salt and  
pepper noise →




← Median  
filtered



Plots of a row of the image

# Median filter

- Median filter preserves sharp transitions (i.e., edges),





... but it removes small brightness variations.


# High-pass filtering (edge detection)

# Edge detection


- Ultimate goal of edge detection: an idealized line drawing.
- Edge contours in the image correspond to important scene contours.



# Edges are sharp intensity changes




# Images as functions $F(x, y)$



Edges look like steep cliffs

# Derivatives and edges

An edge is a place of rapid change in the image intensity function.





# Differentiation and convolution

For a 2D function  $F(x, y)$  the partial derivative along  $x$  is:

$$\frac{\partial F(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{F(x + \varepsilon, y) - F(x, y)}{\varepsilon}$$


For discrete data, we can approximate using finite differences:

$$\frac{\partial F(x, y)}{\partial x} \approx \frac{F(x + 1, y) - F(x, y)}{1}$$

What would be the respective filters along  $x$  and  $y$  to implement the partial derivatives as a convolution?



# Partial derivatives of an image

$$\frac{\partial F(x, y)}{\partial x}$$



$$\frac{\partial F(x, y)}{\partial y}$$

-1	1
----	---



-1
1


# Alternative Finite-difference filters

Prewitt filter  $G_x = \begin{bmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{bmatrix}$  and  $G_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ +1 & +1 & +1 \end{bmatrix}$

Sobel filter  $G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$  and  $G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$

Sample Matlab code




```
>> im = imread('lion.jpg');  
>> h = fspecial('sobel');  
>> outim = imfilter(double(im), h);  
>> imagesc(outim);  
>> colormap gray;
```



# Image gradient

The gradient of an image:  $\nabla F = \left[ \frac{\partial F}{\partial x}, \frac{\partial F}{\partial y} \right]$

The gradient points in the direction of fastest intensity change



The gradient direction (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left( \frac{\frac{\partial F}{\partial y}}{\frac{\partial F}{\partial x}} \right)$$


The *edge strength* is given by the gradient magnitude

$$\|\nabla F\| = \sqrt{\left( \frac{\partial F}{\partial x} \right)^2 + \left( \frac{\partial F}{\partial y} \right)^2}$$

# Effects of noise


Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

# Solution: smooth first




Where is the edge?

Look for peaks in  $\frac{\partial}{\partial x}(F * H)$

# Alternative: combine derivative and smoothing filter

$$\frac{\partial}{\partial x}(F * H) = F * \frac{\partial H}{\partial x}$$


Differentiation property of convolution.



# Derivative of Gaussian filter (along x)


$$(I * G) * H = I * (G * H)$$

$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} * \begin{bmatrix} 1 & -1 \end{bmatrix}$$








# Derivative of Gaussian filters



$x$ -direction




$y$ -direction



# Laplacian of Gaussian


Consider 
$$\frac{\partial^2}{\partial x^2} (F * H) = F * \frac{\partial^2 H}{\partial x^2}$$



Where is the edge?


Zero-crossings of bottom graph

# 2D edge detection filters




Gaussian

$$G = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial G}{\partial x}$$



Laplacian of Gaussian

$$\nabla^2 G$$

- $\nabla^2$  is the Laplacian operator:  $\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$

# Summary on (linear) filters

- Smoothing filter:
  - **has positive values**
  - **sums to 1** → preserve brightness of constant regions
  - removes “high-frequency” components: “low-pass” filter
- Derivative filter:
  - **has opposite signs** used to get high response in regions of high contrast
  - **sums to 0** → no response in constant regions
  - **highlights “high-frequency”** components: “high-pass” filter

# The Canny edge-detection algorithm (1986)

- Compute gradient of smoothed image in both directions
- Discard pixels whose gradient magnitude is below a certain threshold
- **Non-maximal suppression:** identify local maxima along gradient direction

# The Canny edge-detection algorithm (1986)



Take a grayscale image. If not grayscale (i.g., RGB), convert it into a grayscale by replacing each pixel by the mean value of its R, G, B components.

Original image (Lenna image: <https://en.wikipedia.org/wiki/Lenna>)

# The Canny edge-detection algorithm (1986)



Take a grayscale image. If not grayscale (i.g., RGB), convert it into a grayscale by replacing each pixel by the mean value of its R, G, B components.

Original image (Lenna image: <https://en.wikipedia.org/wiki/Lenna>)

# The Canny edge-detection algorithm (1986)




Convolve the image  $F$  with  $x$  and  $y$  derivatives of Gaussian filter


$$\frac{\partial F}{\partial x} = F * \frac{\partial G}{\partial x}$$

$$\frac{\partial F}{\partial y} = F * \frac{\partial G}{\partial y}$$

$$\frac{\partial G}{\partial x}$$



$$\frac{\partial G}{\partial y}$$



$$\|\nabla F\| = \sqrt{\left(\frac{\partial F}{\partial x}\right)^2 + \left(\frac{\partial F}{\partial y}\right)^2} : \text{Edge strength}$$



# The Canny edge-detection algorithm (1986)



Threshold it (i.e., set to 0 all pixels whose value is below a given threshold)

Thresholded  $\|\nabla F\|$

# The Canny edge-detection algorithm (1986)



Take local maximum  
along gradient direction

Thinning: non-maxima suppression (local-maxima detection)  
along edge direction

# Summary (things to remember)

- Image filtering (definition, motivation, applications)
- Moving average
- Linear filters and formulation: box filter, Gaussian filter
- Boundary issues
- Non-linear filters
  - Median filter and its applications
- Edge detection
  - Derivating filters (Prewitt, Sobel)
  - Combined derivative and smoothing filters (deriv. of Gaussian)
  - Laplacian of Gaussian
  - Canny edge detector
- Readings: Ch. 3.2, 4.2.1 of Szeliski book

# Understanding Check

Are you able to:

- Explain the differences between convolution and correlation?
- Explain the differences between a box filter and a Gaussian filter?
- Explain why should one increase the size of the kernel of a Gaussian filter if  $\sigma$  is large (i.e. close to the size of the filter kernel)?
- Explain when would we need a median filter?
- Explain how to handle boundary issues?
- Explain the working principle of edge detection with a  $1D$  signal?
- Explain how noise does affect this procedure?
- Explain the differential property of convolution?
- Show how to compute the first derivative of an image intensity function along  $x$  and  $y$ ?
- Explain why the Laplacian of Gaussian operator is useful?
- List the properties of smoothing and derivative filters?
- Illustrate the Canny edge detection algorithm?
- Explain what non-maxima suppression is and how it is implemented?