University of Zurich[UZH] | ETH zürich

Institute of Informatics – Institute of Neuroinformatics

ROBOTICS & PERCEPTION GROUP

Lecture 07
Multiple View Geometry 1

Davide Scaramuzza

http://rpg.ifi.uzh.ch/

Lab Exercise 5 - Today afternoon

- Room ETH HG E 1.1 from 13:15 to 15:00
- Work description: Stereo vision: rectification, epipolar matching, disparity, triangulation

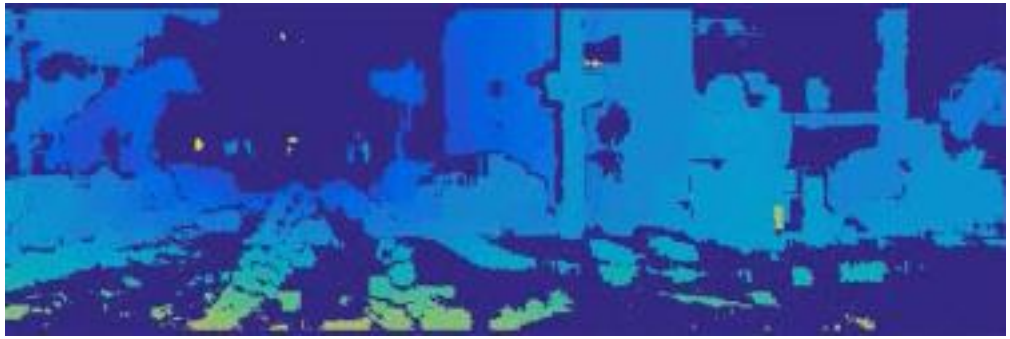

Disparity map

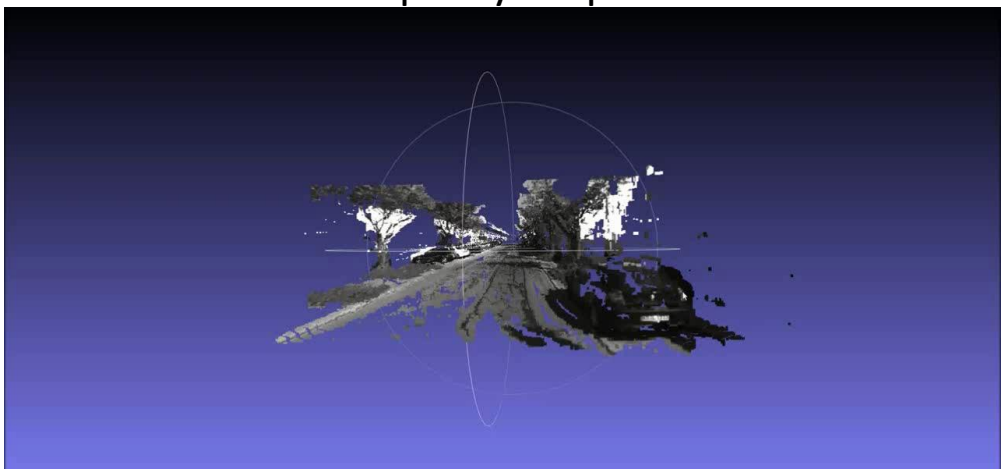

3D point cloud

Course Topics

- Principles of image formation
- Image Filtering
- Feature detection and matching
- Multi-view geometry
- Visual place recognition
- Event-based Vision
- Dense reconstruction
- Visual inertial fusion

Multiple View Geometry

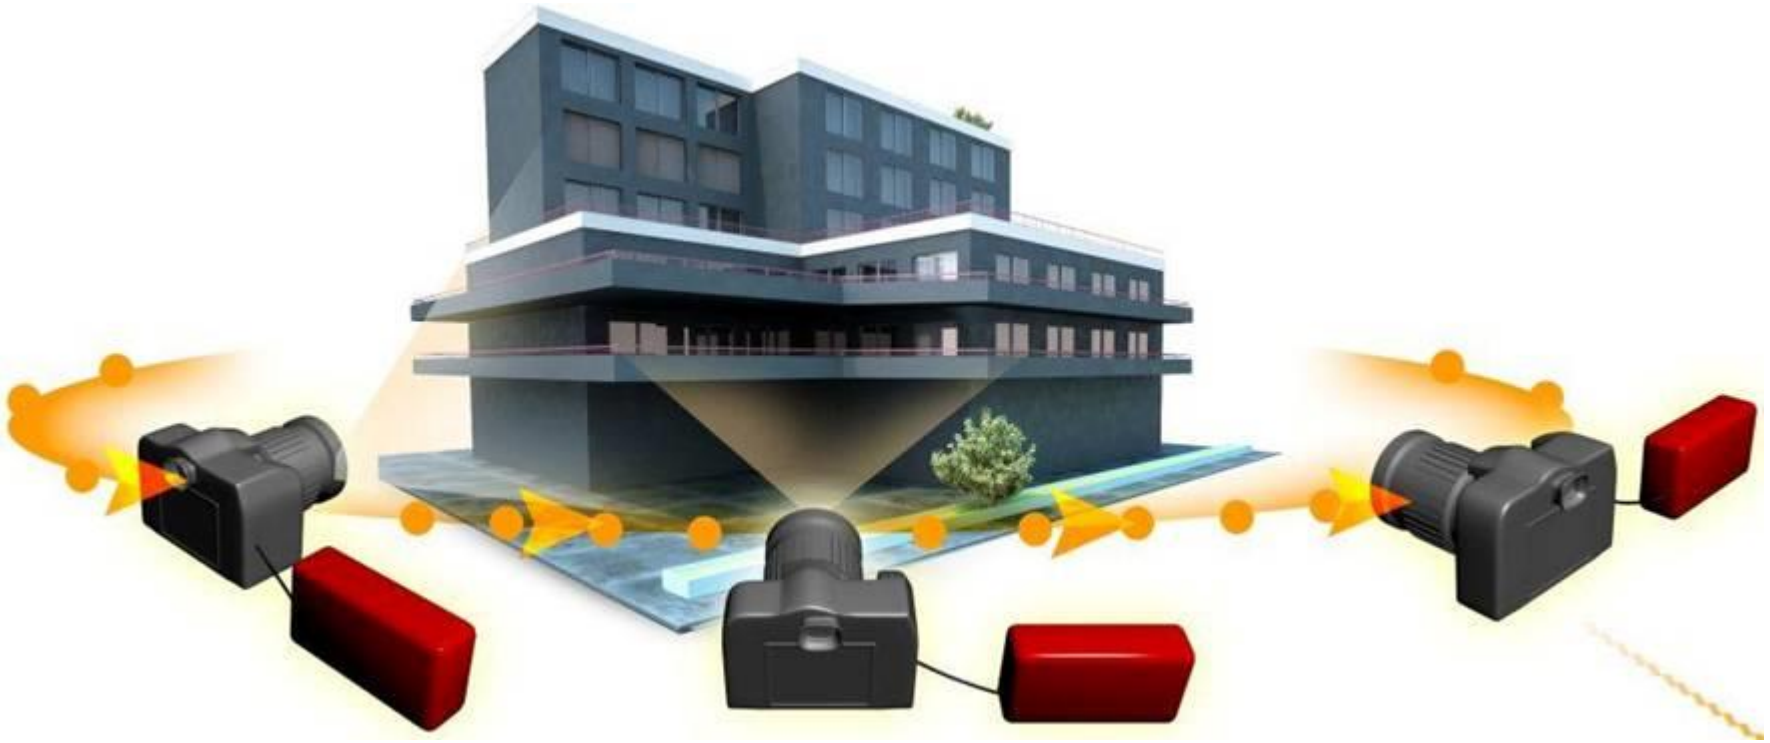

Multiple View Geometry

San Marco square, Venice
14,079 images, 4,515,157 points

Multiple View Geometry

- **3D reconstruction from multiple views:**
  - **Assumptions:** K, T and R are known.
  - **Goal**: Recover the 3D structure from images

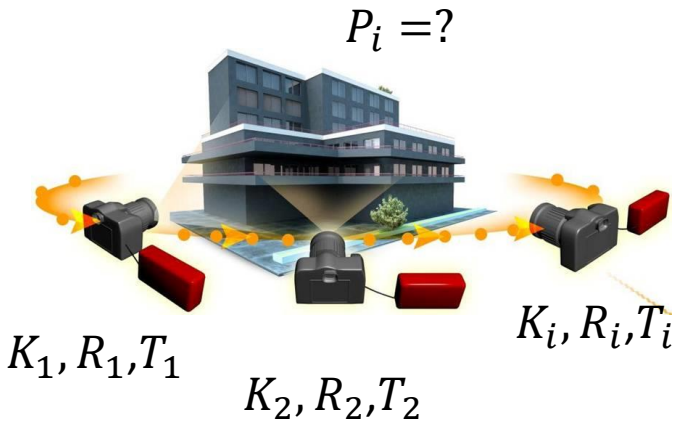

- **Structure From Motion:**
  - **Assumptions**: none (K, T, and R are unknown).
  - **Goal**: Recover simultaneously 3D scene structure and camera poses (up to scale) from multiple images

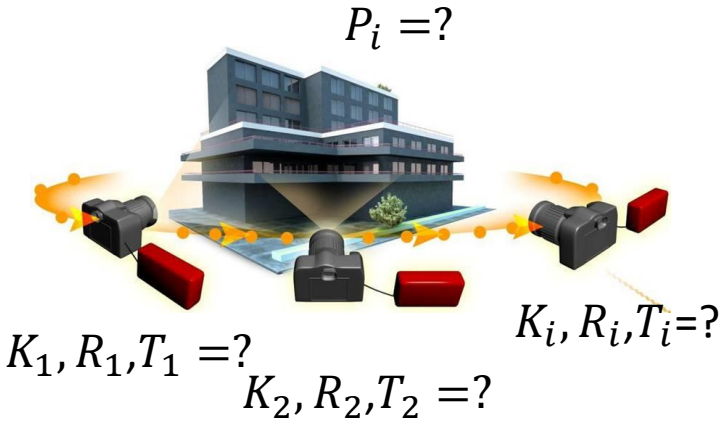

2-View Geometry

- **Depth from stereo (i.e., stereo vision)**
  - **Assumptions:** K, T and R are known.
  - **Goal**: Recover the 3D structure from images

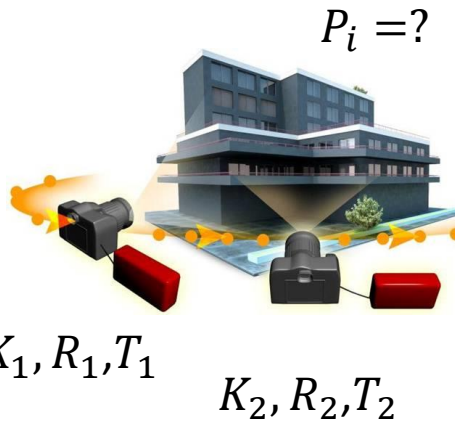

- **2-view Structure From Motion:**
  - **Assumptions**: none (K, T, and R are unknown).
  - **Goal**: Recover simultaneously 3D scene structure, camera poses (up to scale), and intrinsic parameters from two different views of the scene

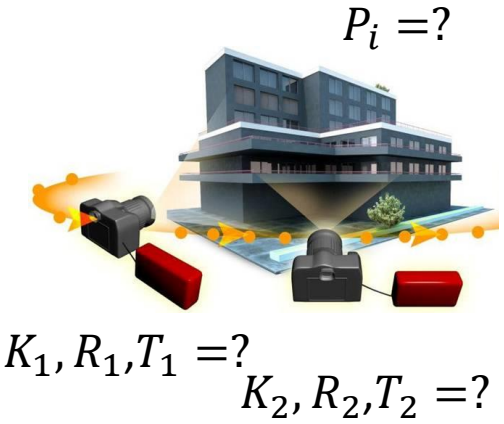

Today's outline

- Stereo Vision
- Epipolar Geometry

Depth from Stereo

- From a single camera, we can only compute the **ray** on which each image point lies
- With a stereo camera (binocular), we can solve for the intersection of the rays and recover the 3D structure

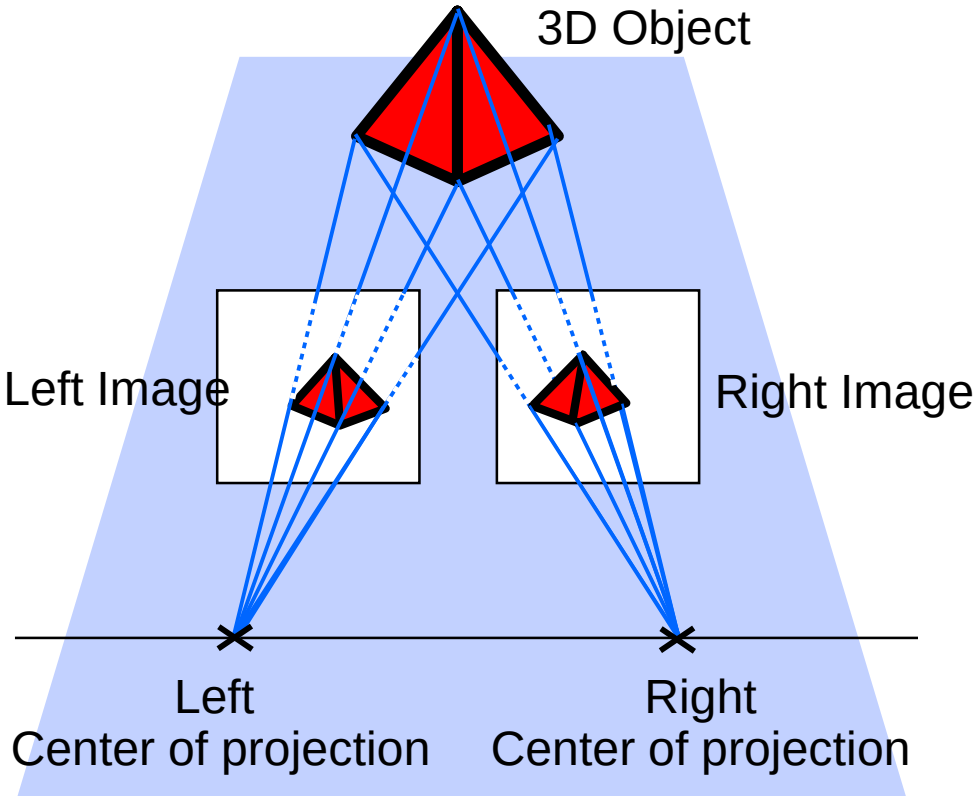

The “human” binocular system

- **Stereopsys:** the brain allows us to see the left and right retinal images as a single 3D image
- The images project on our retina up-side-down but our brains lets us perceive them as «straight». Radial distortion is also removed. This process is called «**rectification**». What happens if you wear a pair of mirrors for a week?

Image from the left eye

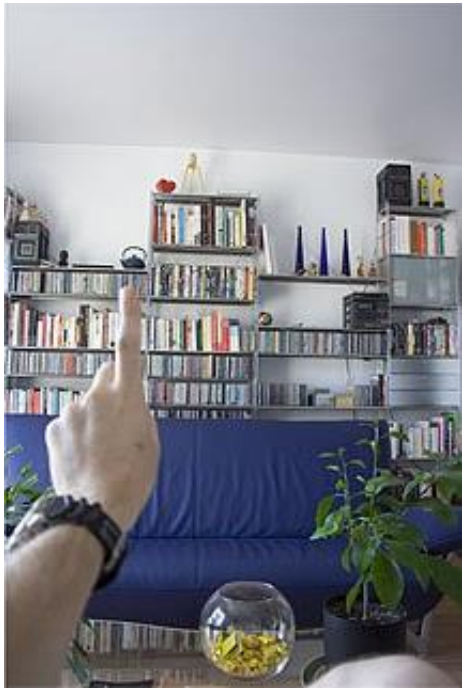

Image from the right eye

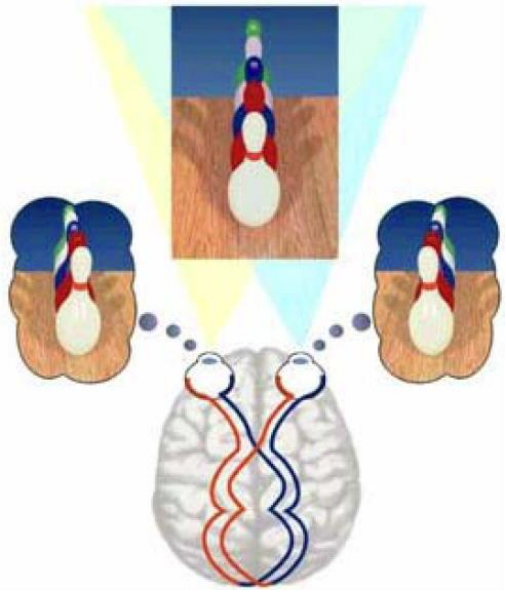

The “human” binocular system

- **Stereopsys:** the brain allows us to see the left and right retinal images as a single 3D image
- The images project on our retina up-side-down but our brain lets us perceive them as «straight». Radial disotion is also removed. This process is called «**rectification**»

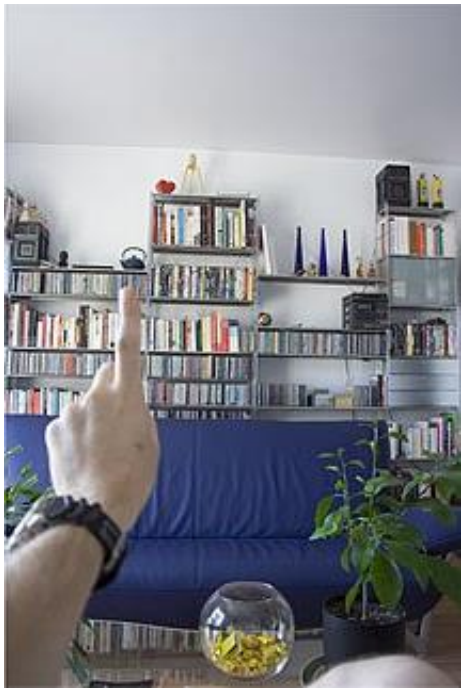

Make a simple test:

1. Fix an object
2. Open and close alternatively the left and right eyes.

- The horizontal displacement is called **disparity**
- The smaller the disparity, the farther the object

The “human” binocular system

- **Stereopsys:** the brain allows us to see the left and right retinal images as a single 3D image
- The images project on our retina up-side-down but our brains lets us perceive them as «straight». Radial disotion is also removed. This process is called «**rectification**»

Make a simple test:

1. Fix an object
2. Open and close alternatively the left and right eyes.

- The horizontal displacement is called **disparity**
- The smaller the disparity, the farther the object

Disparity

- The disparity between the left and right image allows us to perceive the depth

These animated GIF images display intermittently the left and right image

Applications: Stereograms

Exploit disparity as depth cue using single image

Image from magiceye.com

Applications: Stereograms

Exploit disparity as depth cue using single image

Image from magiceye.com

Applications: Stereo photography and stereo viewers

Take two pictures of the same subject from two different viewpoints and display them so that each eye sees only one of the images.

Invented by Sir Charles Wheatstone, 1838

Applications: Anaglyphs

The first method to produce anaglyph images was developed in 1852 by Wilhelm Rollmann in Leipzig, Germany

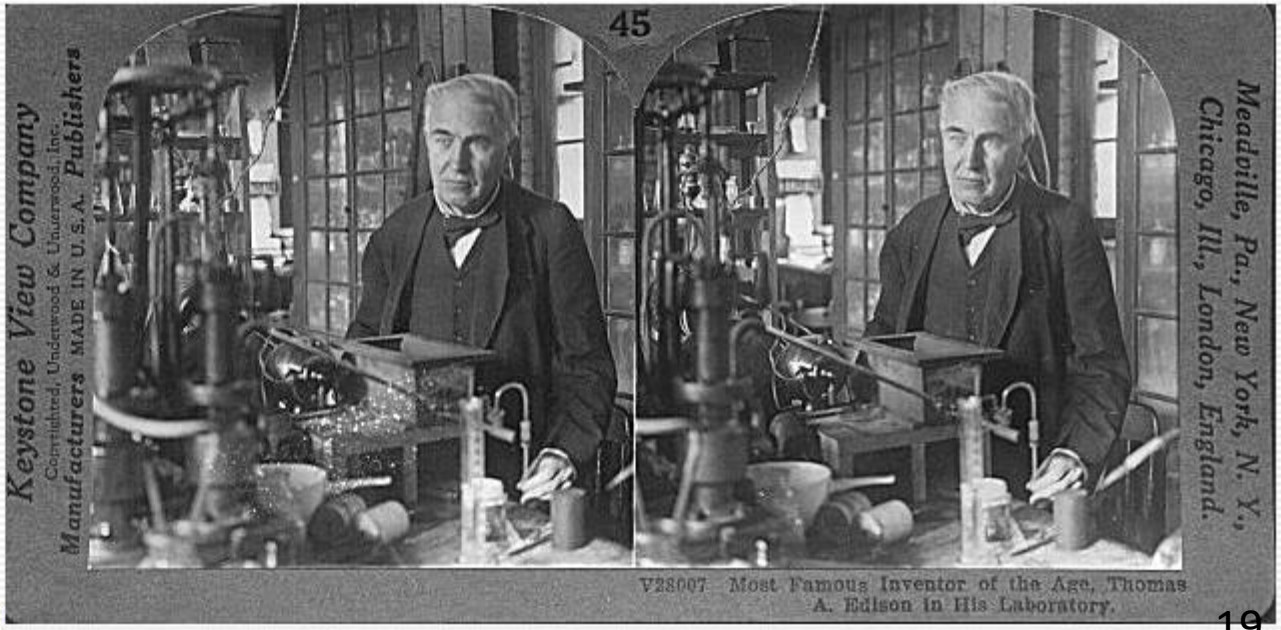

Stereo Vision

- Triangulation
  - Simplified case
  - General case
- Correspondence problem
- Stereo rectification

Stereo Vision: basic idea

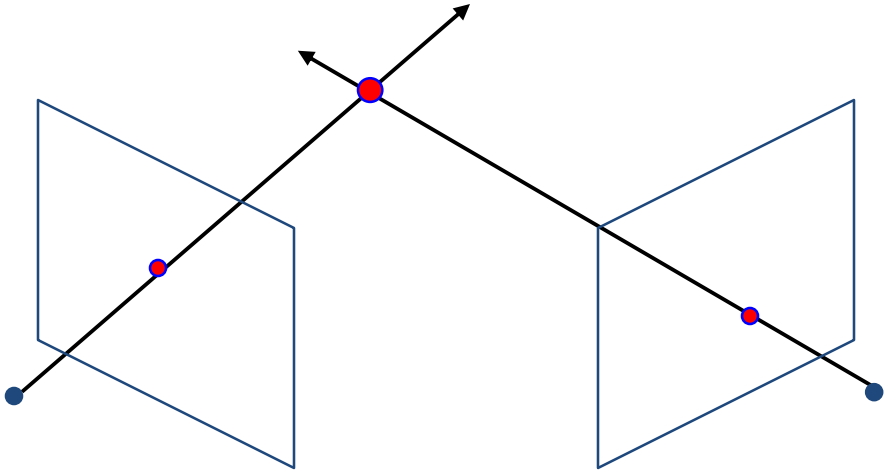

Basic Principle: Triangulation

- Gives reconstruction as intersection of two rays
- Requires
  - camera pose (calibration)
  - point correspondence

Stereo Vision: basic idea

Simplified case

(identical cameras and aligned)

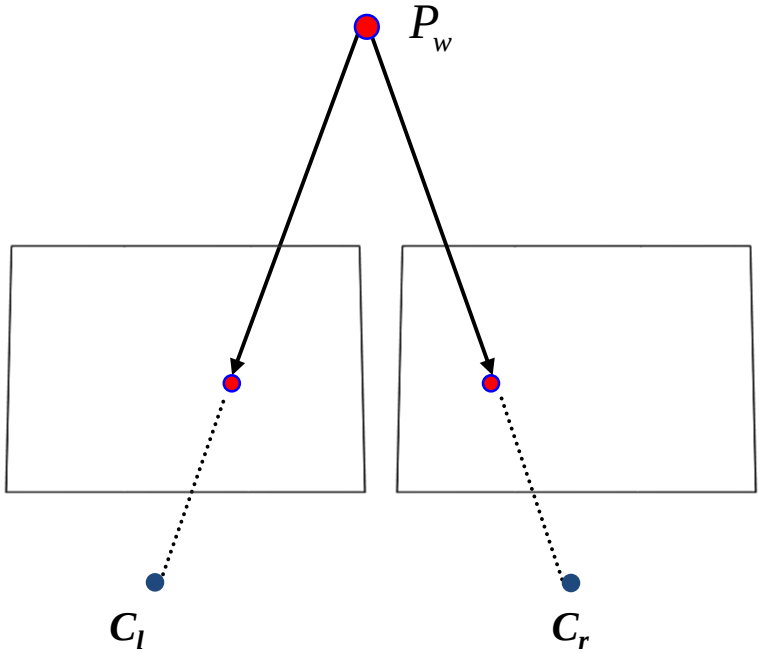

General case

(non identical cameras and not aligned)

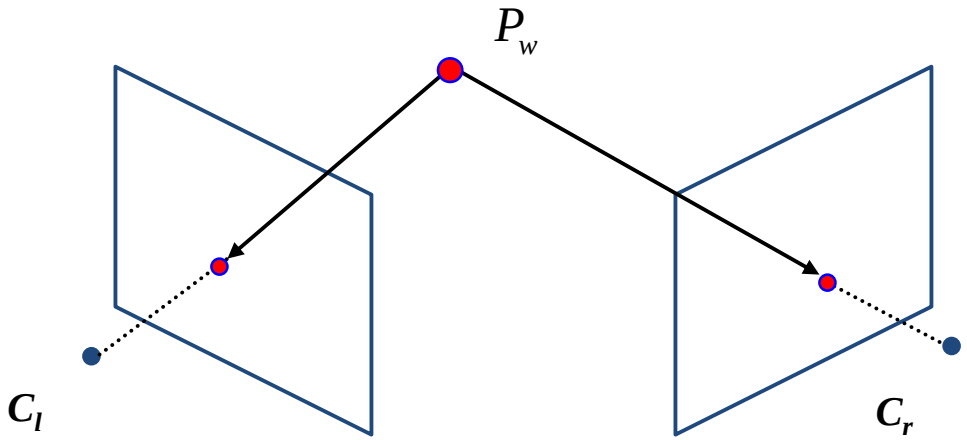

Stereo Vision - The simplified case

Both cameras are **identical** (i.e., same focal length) and are **aligned** with the x-axis

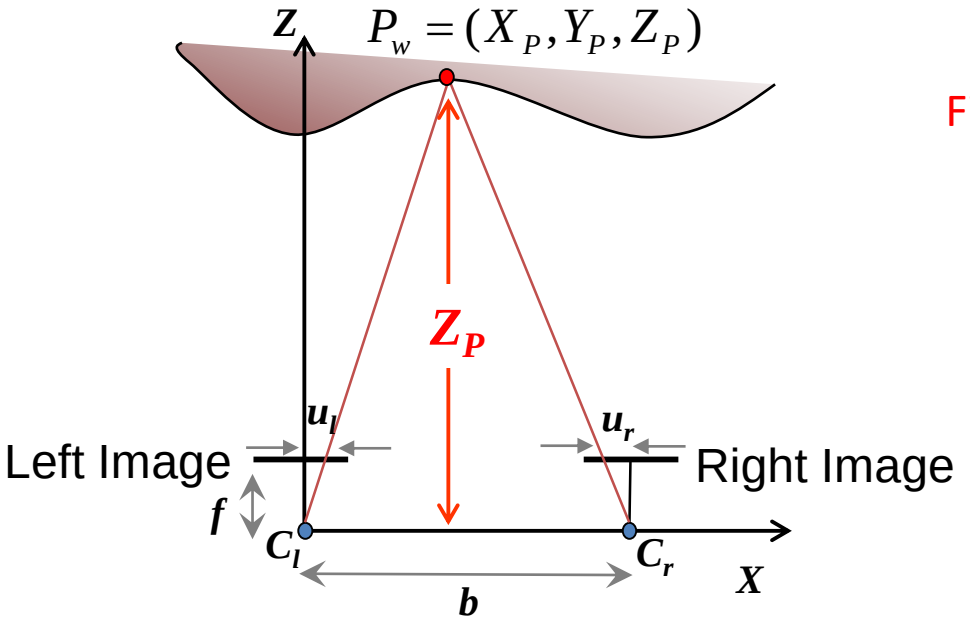

Baseline

distance between the optical centers of the two cameras

Find an expression for the depth Z_P of point P_w

Stereo Vision - The simplified case

Both cameras are **identical** and are **aligned** with the x-axis

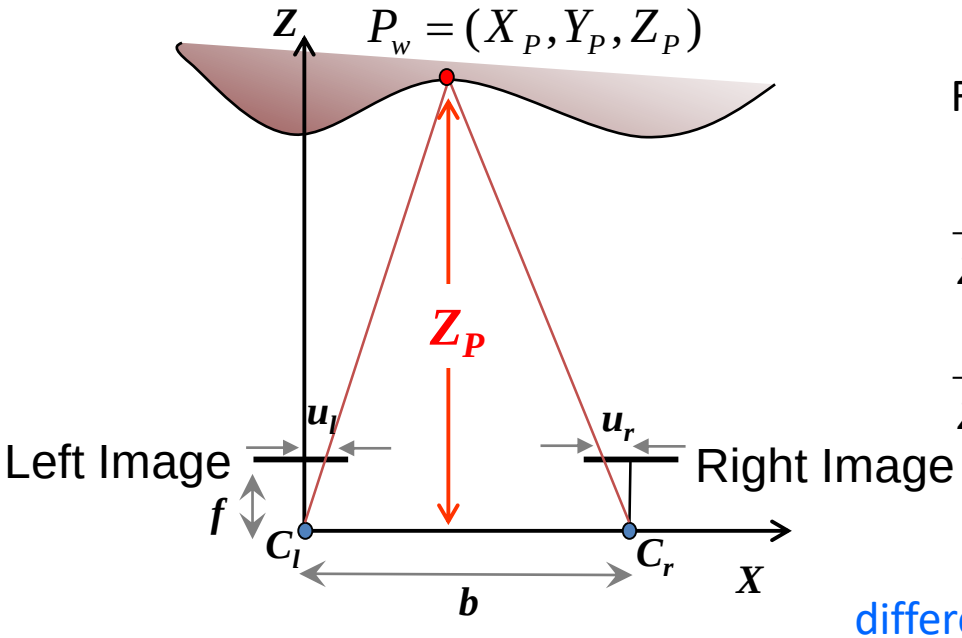

Baseline

distance between the optical centers of the two cameras

From Similar Triangles:

$$\frac{f}{Z_P} = \frac{u_l}{X_P}$$

$$\frac{f}{Z_P} = \frac{-u_r}{b - X_P}$$

$$Z_P = \frac{bf}{u_l - u_r}$$

Disparity

difference in image location of the projection of a 3D point on two image planes

1. What's the max disparity of a stereo camera?
2. What's the disparity of a point at infinity?
3. How does the uncertainty of depth depend on the disparity?
4. And on the depth estimate?
5. How do I increase the accuracy of a stereo system?

Choosing the Baseline

- What's the optimal baseline?
  - **Too small:**
    - Large depth error
    - Can you quantify the error as a function of the disparity?
  - **Too large:**
    - Minimum measurable distance increases
    - Difficult search problem for close objects

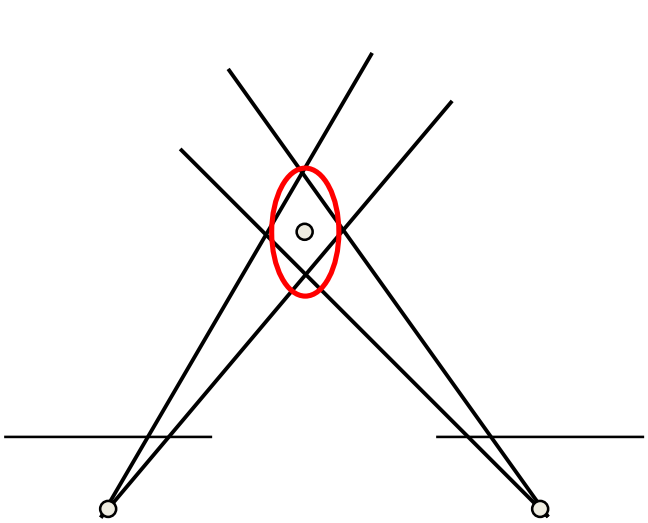

Large Baseline

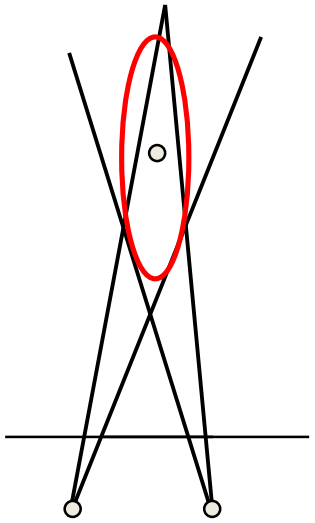

Small Baseline

Stereo Vision – the general case

- Two identical cameras do not exist in nature!
- Aligning both cameras on a horizontal axis is very hard -> Impossible, why?

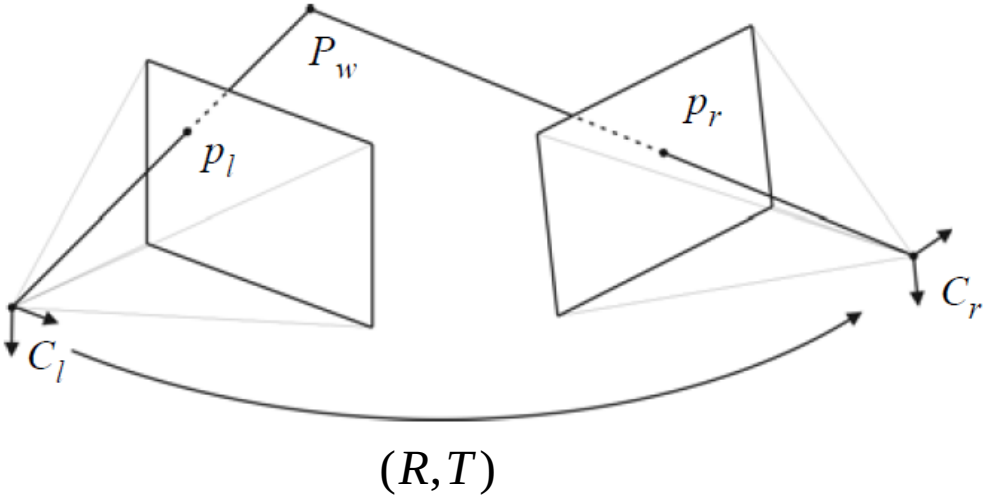

- In order to be able to use a stereo camera, we need the
  - **Extrinsic parameters** (relative rotation and translation)
  - **Instrinsic parameters** (focal length, optical center, radial distortion of each camera)

  ⇨ Use a calibration method (Tsai or Homographies, see Lectures 2, 3)

  ⇨ How do we compute the relative pose?

Stereo Vision – the general case

- To estimate the 3D position of P_w we construct the system of equations of the left and right cameras, and solve it. Do lines always intersect in 3D space?

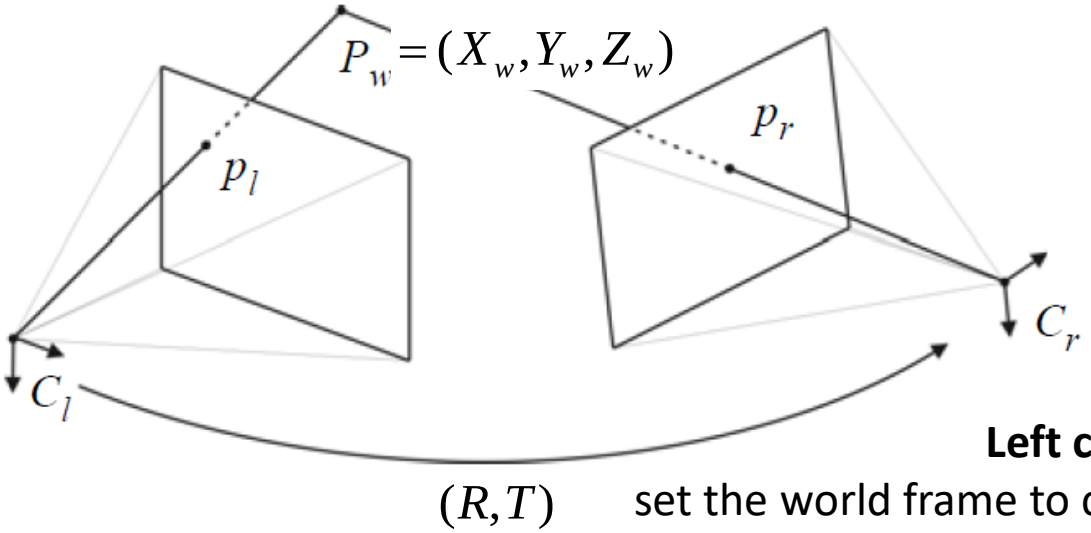

Left camera: set the world frame to coincide with the left camera frame

$$\tilde{p}_l = \lambda_l \begin{bmatrix} u_l \\ v_l \\ 1 \end{bmatrix} = K_l \begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix}$$

Right camera:

$$\tilde{p}_r = \lambda_r \begin{bmatrix} u_r \\ v_r \\ 1 \end{bmatrix} = K_r R \begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} + T$$

- "Triangulation": the problem of determining the 3D position of a point given a set of corresponding image locations and known camera poses.

Review: Cross Product (or Vector Product)

$$\vec{a} \times \vec{b} = \vec{c}$$

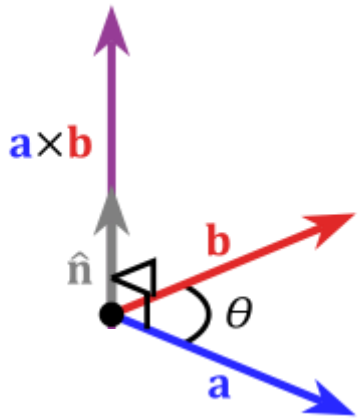

- Vector cross product takes two vectors and returns a third vector that is perpendicular to both inputs

$$\vec{a} \cdot \vec{c} = 0$$

$$\vec{b} \cdot \vec{c} = 0$$

- So here, $\boldsymbol{c}$ is perpendicular to both $\boldsymbol{a}$ and $\boldsymbol{b}$, which means the dot product = 0
- Also, recall that the cross product of two parallel vectors = 0

- The vector **cross product** can also be expressed as the product of a **skew-symmetric matrix** and a vector

$$\mathbf{a} \times \mathbf{b} = \begin{bmatrix} 0 & -a_z & a_y \\ a_z & 0 & -a_x \\ -a_y & a_x & 0 \end{bmatrix} \begin{bmatrix} b_x \\ b_y \\ b_z \end{bmatrix} = [\mathbf{a}_\times]\mathbf{b}$$

Triangulation: Linear Approach

Left camera

$$\lambda_1 \begin{bmatrix} u_1 \\ v_1 \\ 1 \end{bmatrix} = K\begin{bmatrix} I|0 \end{bmatrix} \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix} \Rightarrow \lambda_1 p_1 = M_1 \cdot P$$

Right camera

$$\lambda_2 \begin{bmatrix} u_2 \\ v_2 \\ 1 \end{bmatrix} = K\begin{bmatrix} R|T \end{bmatrix} \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix} \Rightarrow \lambda_2 p_2 = M_2 \cdot P$$

Triangulation: Linear Approach

Left camera

$$\Rightarrow \lambda_1 p_1 = M_1 \cdot P \qquad \Rightarrow p_1 \times M_1 \cdot P = 0 \qquad \Rightarrow \left[p_{1\times}\right] M_1 \cdot P = 0$$

Right camera

$$\Rightarrow \lambda_2 p_2 = M_2 \cdot P \qquad \Rightarrow p_2 \times M_2 \cdot P = 0 \qquad \Rightarrow \left[p_{2\times}\right] M_2 \cdot P = 0$$

Cross product as matrix multiplication:

$$\mathbf{a} \times \mathbf{b} = \begin{bmatrix} 0 & -a_z & a_y \\ a_z & 0 & -a_x \\ -a_y & a_x & 0 \end{bmatrix} \begin{bmatrix} b_x \\ b_y \\ b_z \end{bmatrix} = [\mathbf{a}_\times]\mathbf{b}$$

Triangulation: Linear Approach

Left camera

$$\Rightarrow \lambda_1 p_1 = M_1 \cdot P \qquad \Rightarrow p_1 \times M_1 \cdot P = 0 \qquad \Rightarrow \left[p_{1\times}\right] M_1 \cdot P = 0$$

Right camera

$$\Rightarrow \lambda_2 p_2 = M_2 \cdot P \qquad \Rightarrow p_2 \times M_2 \cdot P = 0 \qquad \Rightarrow \left[p_{2\times}\right] M_2 \cdot P = 0$$

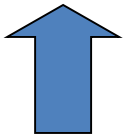

Two independent equations each in terms of three unknown entries of $\boldsymbol{P}$.
$\boldsymbol{P}$ can be determined using SVD, as we already did when we talked about DLT

Triangulation: geometric interpretation

- Given the projections $\mathbf{p_1}$ and $\mathbf{p_2}$ of a 3D point ***P*** in two or more images (with known camera matrices R and T), find the coordinates of the 3D point

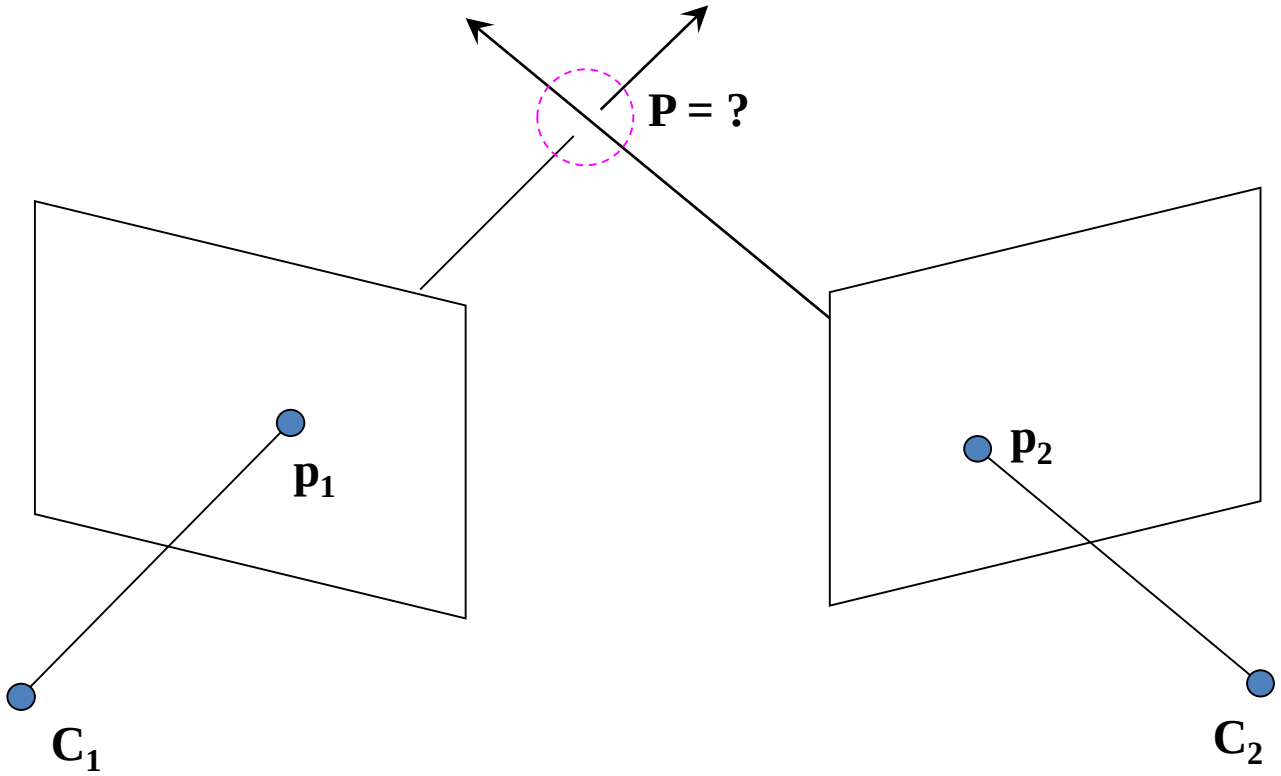

Triangulation: geometric interpretation

- We want to intersect the two visual rays corresponding to $\mathbf{p}_1$ and $\mathbf{p}_2$, but because of noise and numerical errors, they don't meet exactly

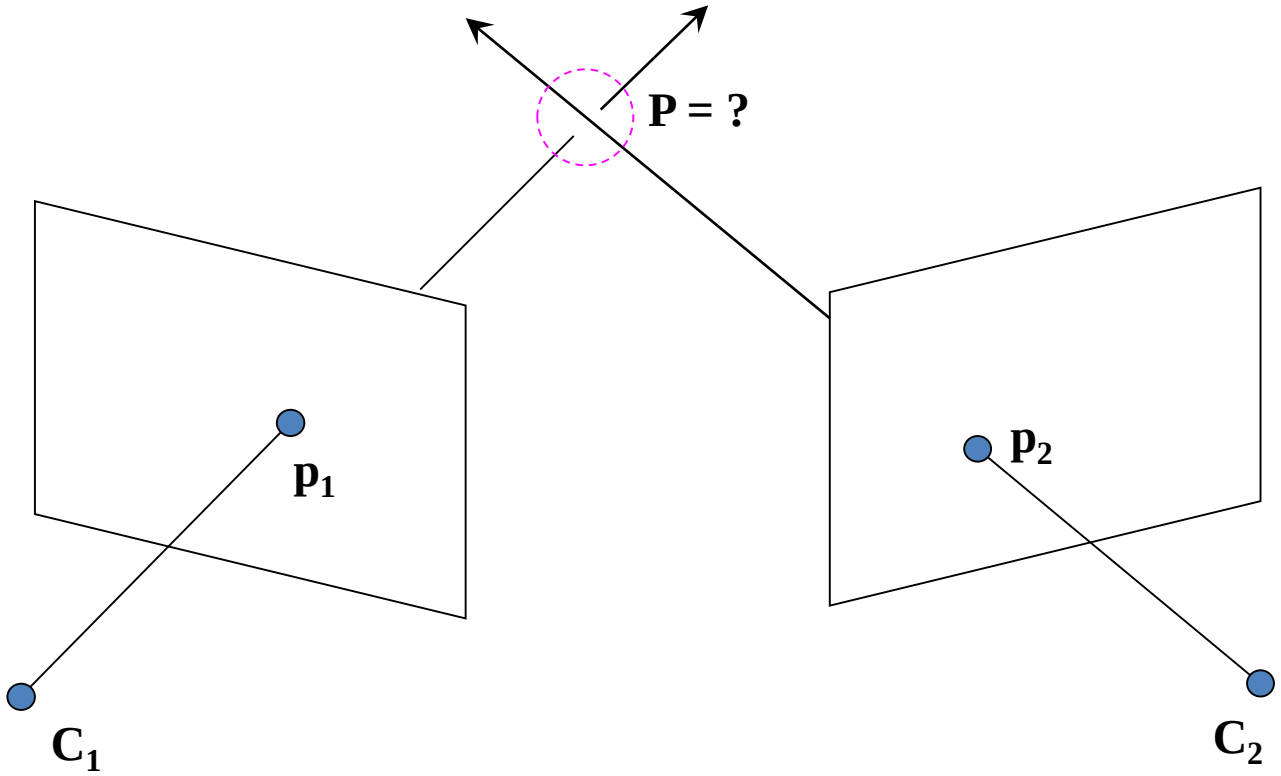

Triangulation: geometric interpretation

- Find shortest segment connecting the two viewing rays and let **P** be the midpoint of that segment

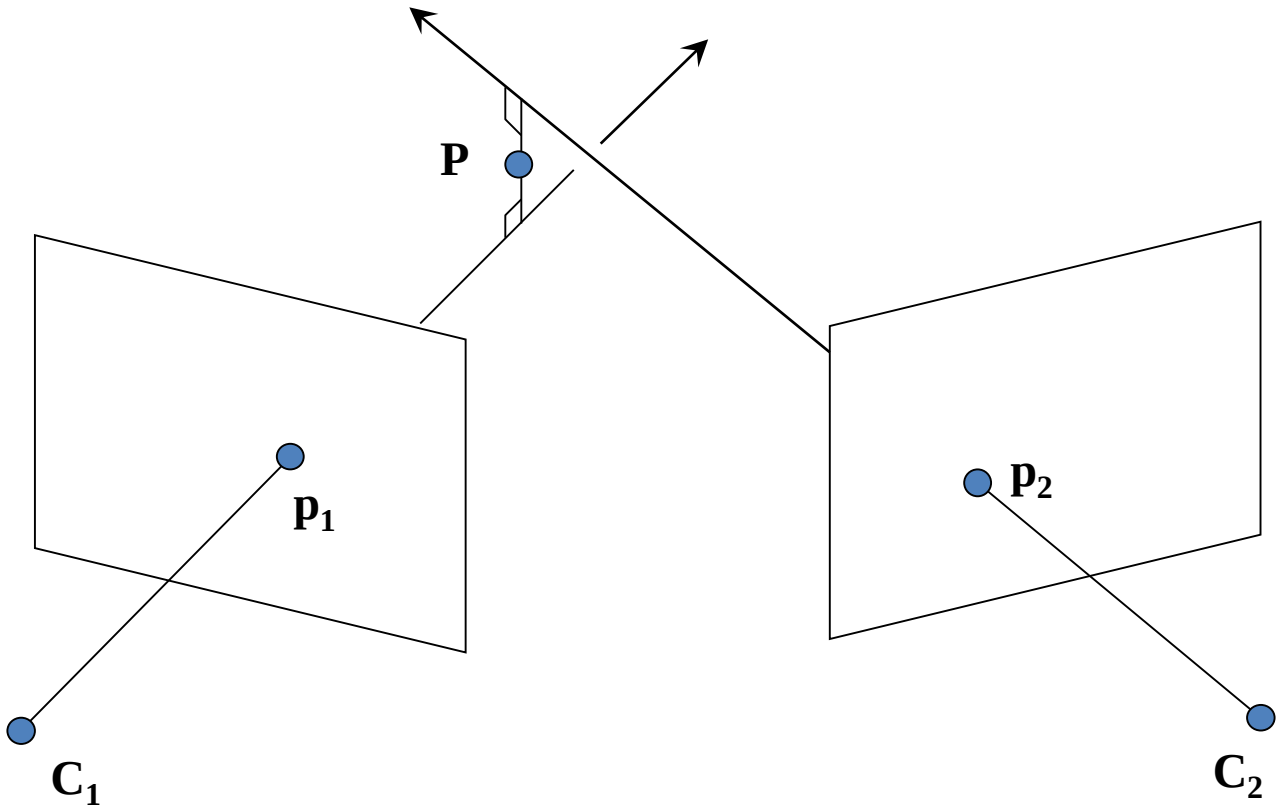

Triangulation: Nonlinear approach

- Find P that minimizes the **Sum of Squared Reprojection Error**:

$$SSRE = d^2(p_1, \pi_1(P)) + d^2(p_2, \pi_2(P))$$

where $d(p_1, \pi_1(P)) = \|p_1 - \pi_1(P)\|$ is called ***Reprojection Error***.

- In practice, initialize ***P*** using linear approach and then minimize SSRE using Gauss-Newton or Levenberg-Marquardt.

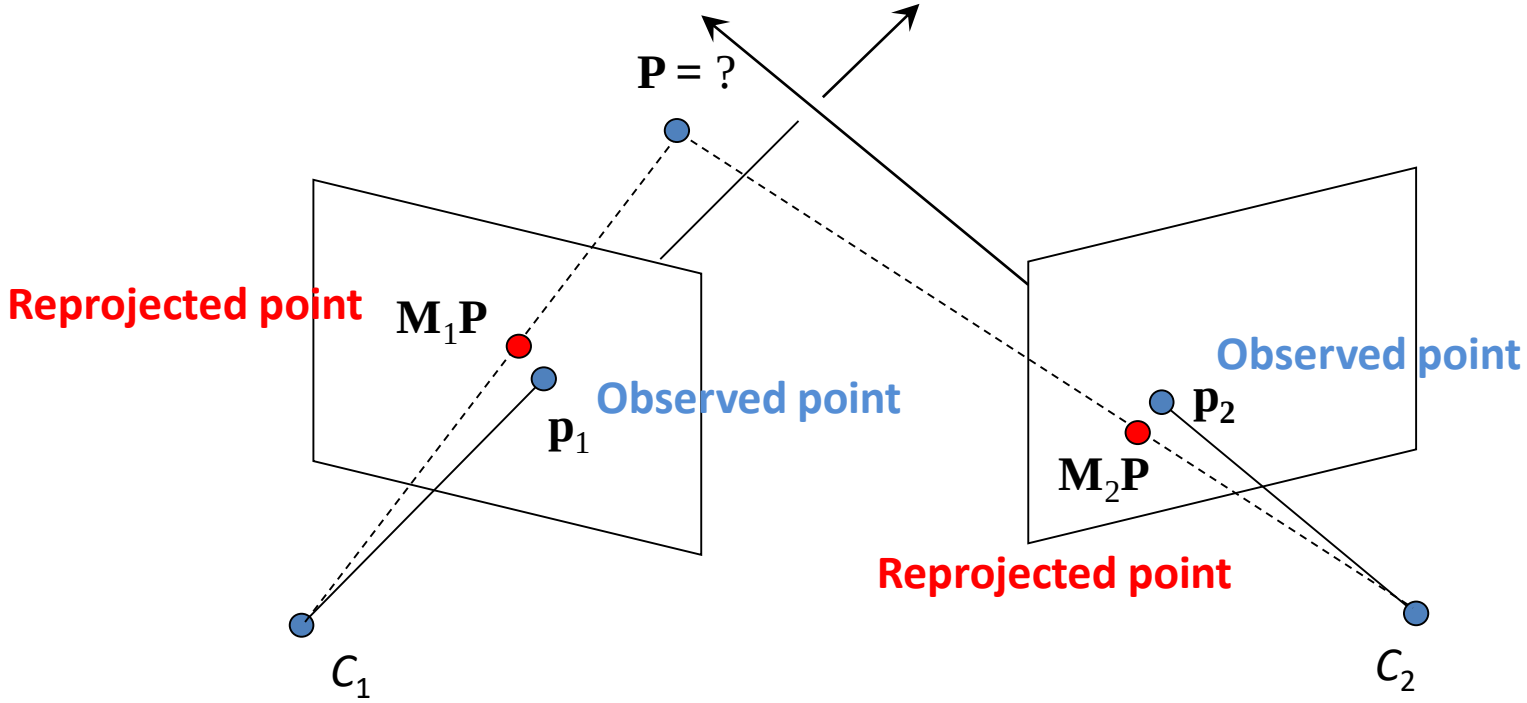

Stereo Vision

- Triangulation
  - Simplified case
  - General case
- Correspondence problem
- Stereo rectification

Correspondence Problem

Given the point p on left image, where is its corresponding point p' on the right image?

Left image

Right image

Correspondence Problem

Given the point p on left image, where is its corresponding point p' on the right image?

Left image

Right image

Correspondence Problem

- **Correspondence search**: identify image patches on the left & right images, corresponding to the same scene structure.
- **Similarity measures**:
  - **(Z)ZNCC**
  - **(Z)SSD**
  - **(Z)SAD**
  - **Census Transform**

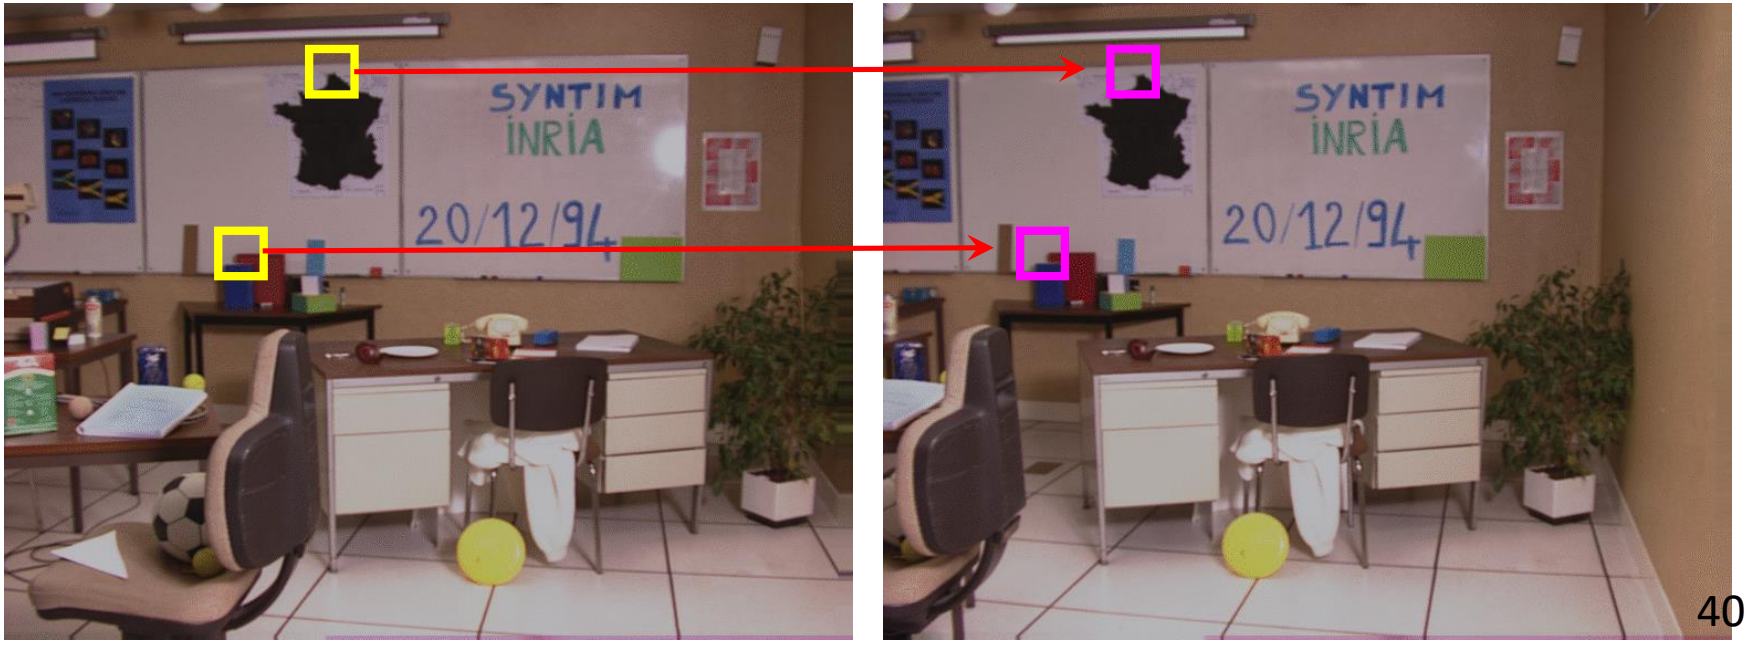

Correspondence Problem

- **Exhaustive** image search can be computationally very expensive!
- Can we make the correspondence search in 1D?
- Potential matches for $\boldsymbol{p}$ have to lie on the corresponding epipolar line $\boldsymbol{l}'$
  - The **epipolar line** is the projection of the infinite ray $\pi^{-1}(p)$ corresponding to $\boldsymbol{p}$ in the other camera image
  - The **epipole** is the projection of the optical center on the other camera image
  - A stereo camera has two epipoles

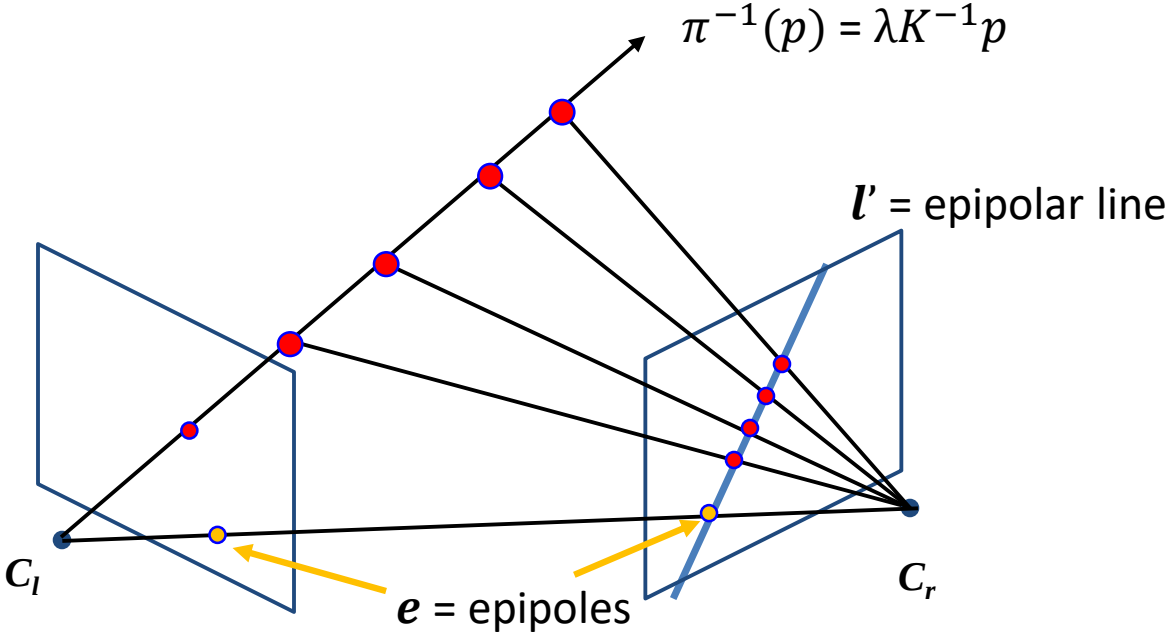

The Epipolar Constraint

- The epipolar plane is uniquely defined by the two optical centers C_l, C_r and one image point p
- The **epipolar constraint** constrains the location, in the second view, of the corresponding point to a given point in the first view.
- Why is this useful?
  - Reduces correspondence problem to 1D search along *conjugate epipolar lines*

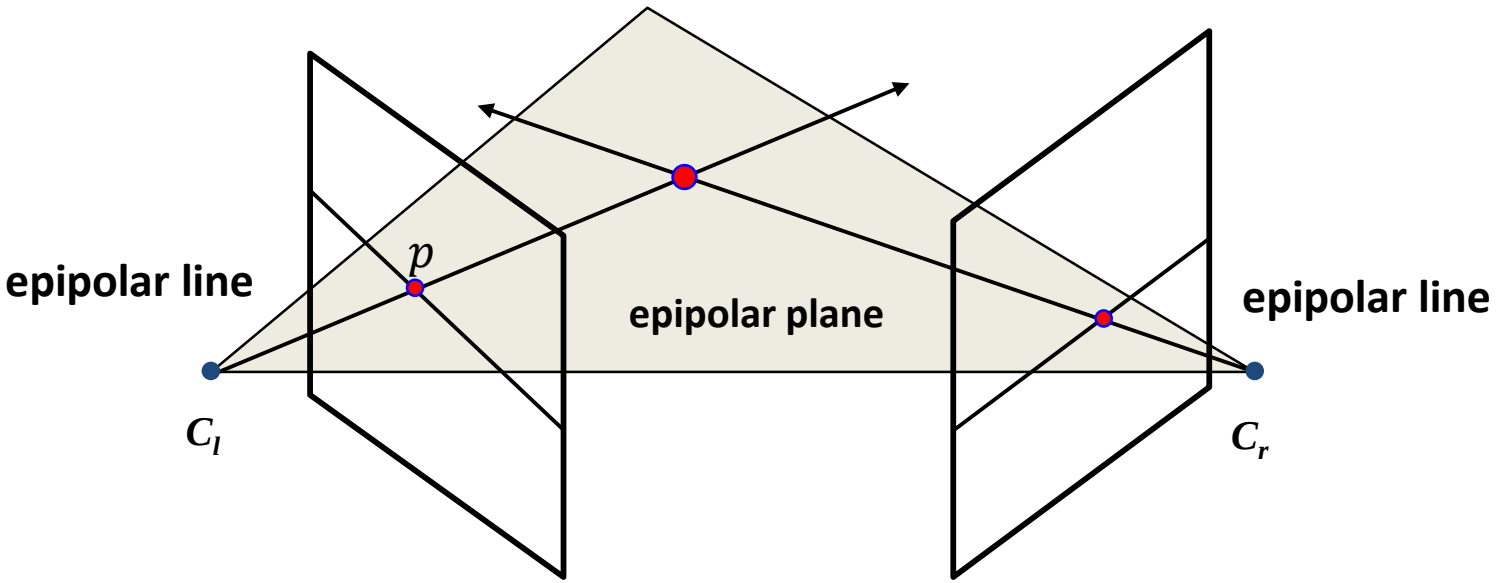

Correspondence Problem: Epipolar Constraint

Thanks to the epipolar constraint, corresponding points can be searched for, along epipolar lines: ⇨ computational cost reduced to 1 dimension!

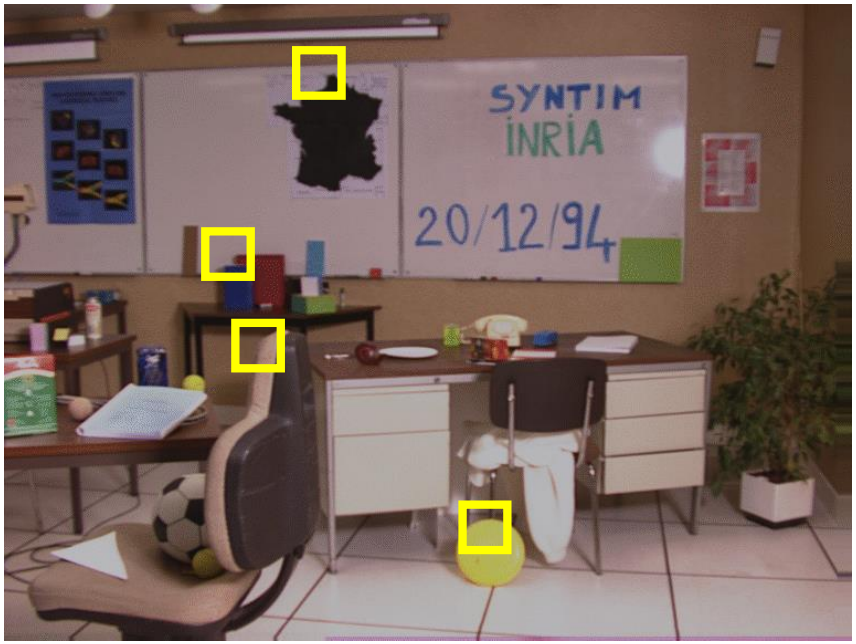

Left image

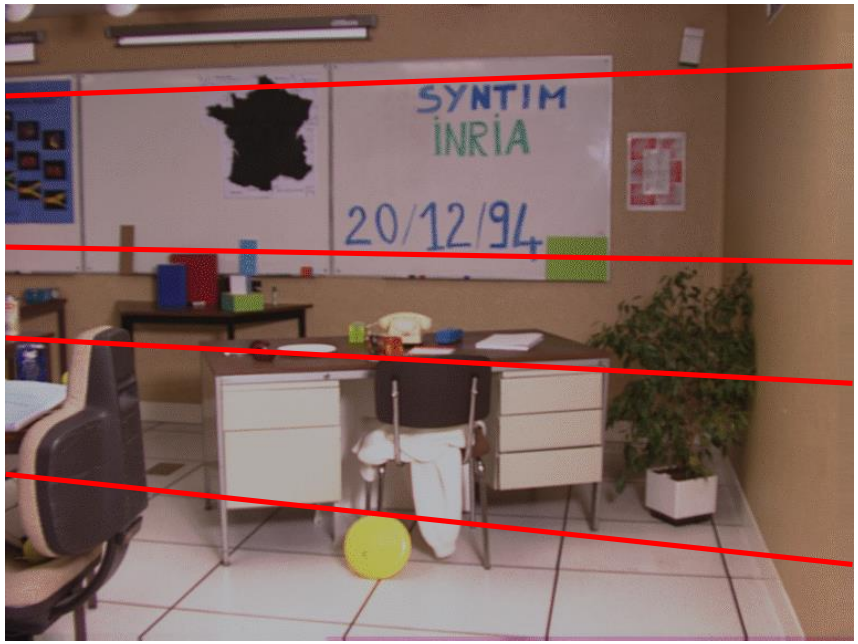

Right image

Example: converging cameras

- **Remember**: all the epipolar lines intersect at the epipole
- As the position of the 3D point varies, the epipolar lines “rotate” about the baseline

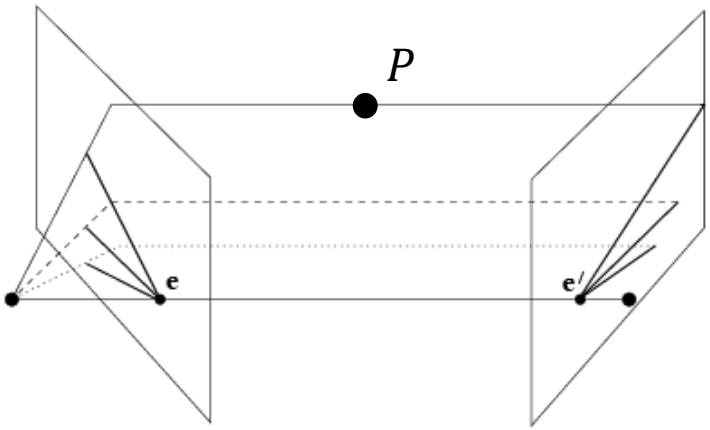

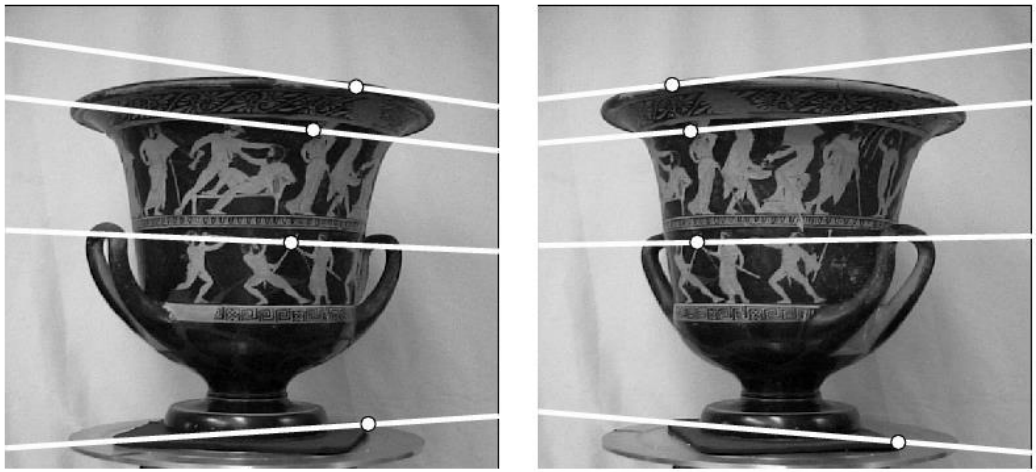

Left image

Right image

Example: identical and horizontally-aligned cameras

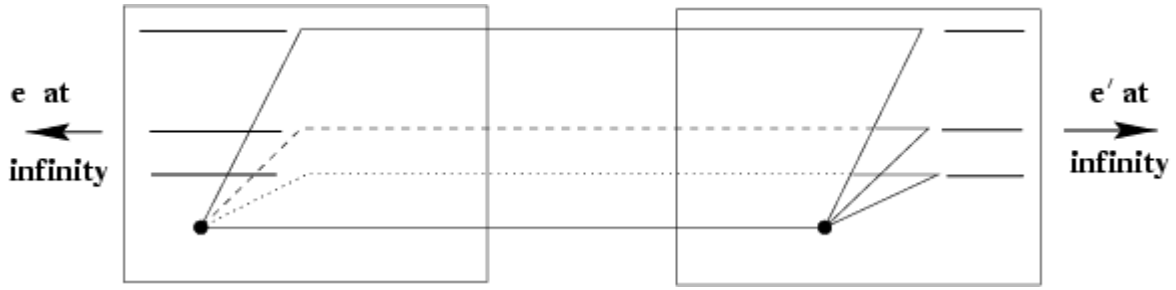

Left image

Right image

Example: forward motion (parallel to the optical axis)

- Epipole has the **same coordinates** in both images
- Points move along lines radiating from e: “Focus of expansion”

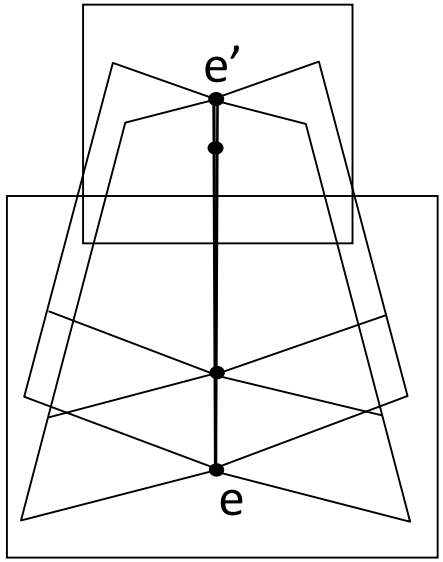

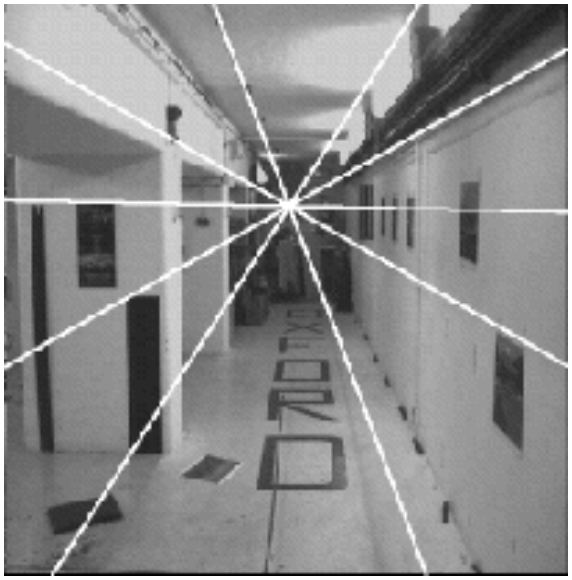

Left image

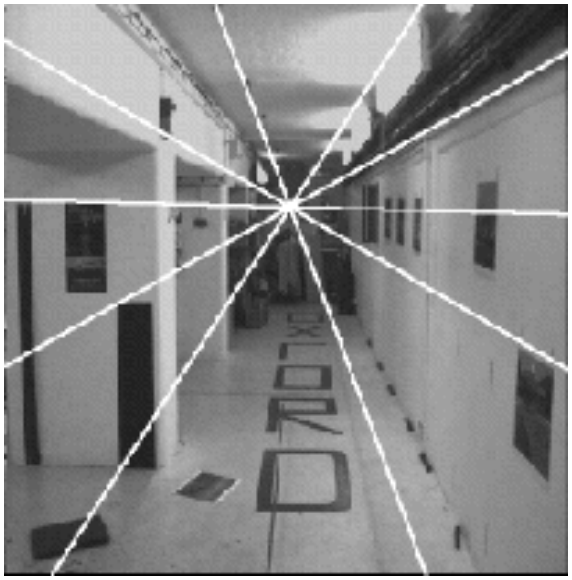

Right image

Stereo Vision

- Simplified case
- General case
- Correspondence problem
- Stereo rectification
- Triangulation

Stereo Rectification

- Even in commercial stereo cameras the left and right images are never perfectly aligned.
- In practice, it is convenient if image scanlines are the epipolar lines.
- Stereo rectification warps the left and right images into new “rectified” images, whose epipolar lines are aligned to the baseline.

Stereo Rectification

- Reprojects image planes onto a common plane parallel to the baseline
- It works by computing two homographies, one for each input image reprojection

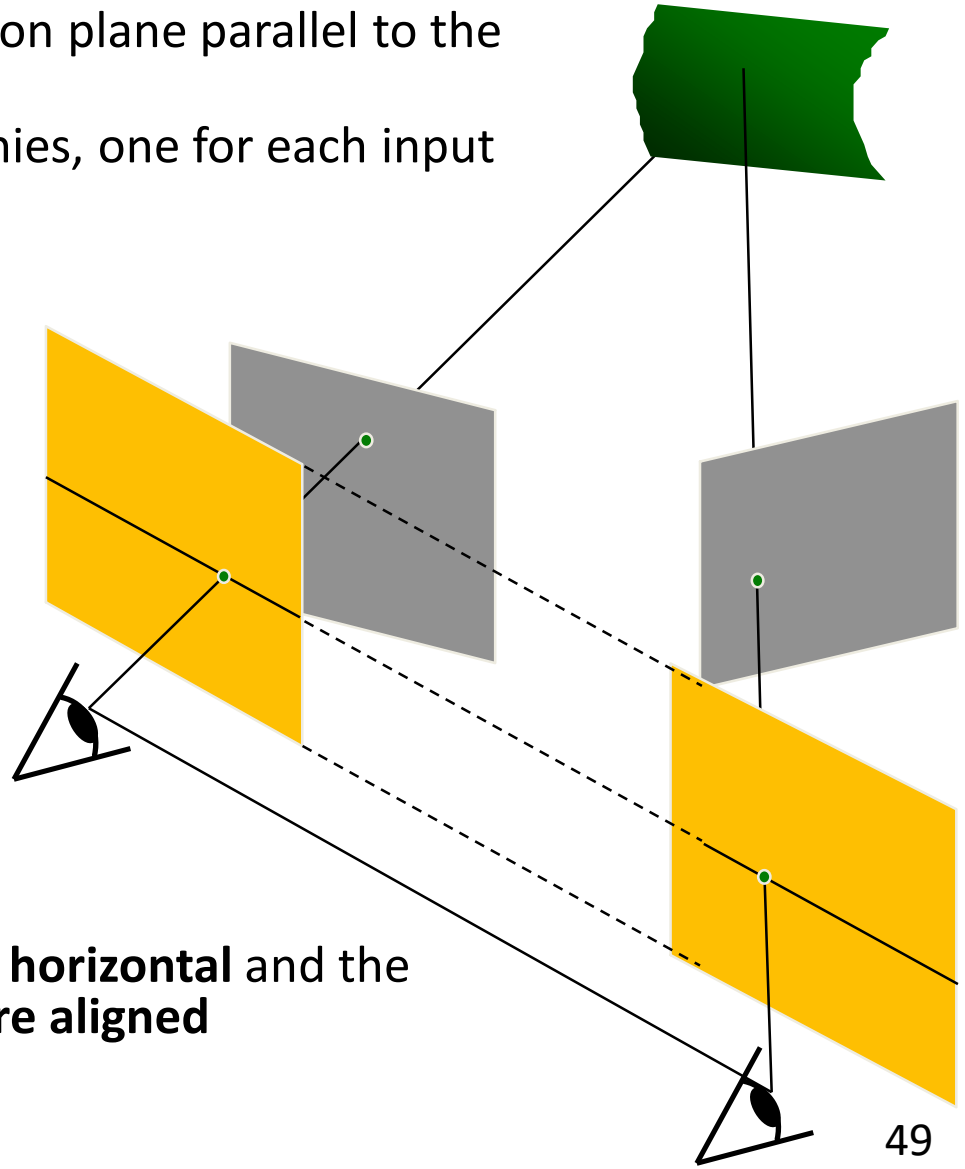

- As a result, the new **epipolar lines** are **horizontal** and the **scanlines** of the left and right image **are aligned**

Stereo Rectification

- The idea behind rectification is to define two new Perspective Projection Matrices (PPMs) obtained by rotating the old ones around their optical centers until focal planes become parallel to each other.
- This ensures that epipoles are at infinity, hence epipolar lines are parallel.
- To have horizontal epipolar lines, the baseline must be parallel to the new X axis of both cameras.
- In addition, to have a proper rectification, corresponding points must have the same vertical coordinate. This is obtained by requiring that the new cameras have the same intrinsic parameters.
- Note that, being the focal length the same, the new image planes are coplanar too
- PPMs are the same as the old cameras, whereas the new orientation (the same for both cameras) differs from the old ones by suitable rotations; intrinsic parameters are the same; for both cameras.

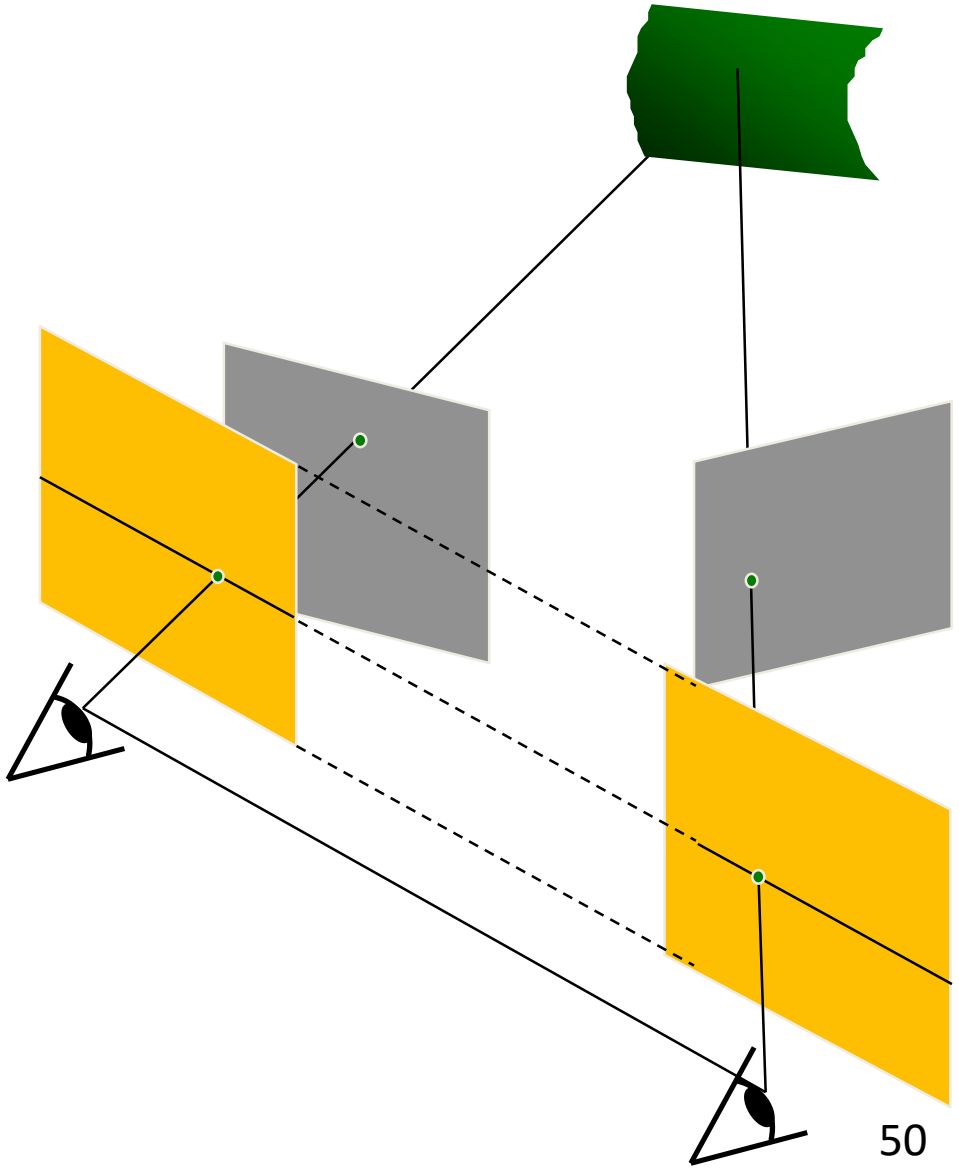

Stereo Rectification (1/5)

We have seen in Lecture 02 that the Perspective Equation for a point P_w in the world frame is

$$\lambda \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K[R|T] \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

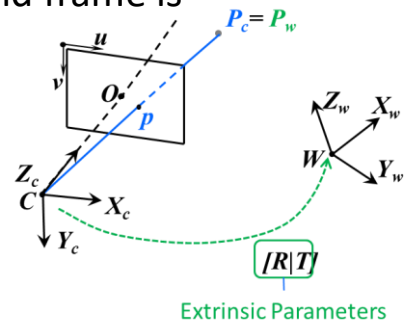

where $R = R_{cw}$ and $T = T_{cw}$ transform points **from the World frame to the Camera frame**.

This can be re-written in a more convenient way. In the following, let $R \equiv R_{wc}$ and $T \equiv T_{wc}$ transform points **from the Camera frame to the World frame**. (Recall that $R_{wc} = R_{cw}^{-1} = R_{cw}^{t}$ and that $T_{wc} = -R_{cw}^{t} T_{cw} = C$ is the optical center of the camera). The projection equation can be re-written as:

$$\lambda \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = KR^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - T \right)$$

$$\Rightarrow \lambda \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = KR^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C \right)$$

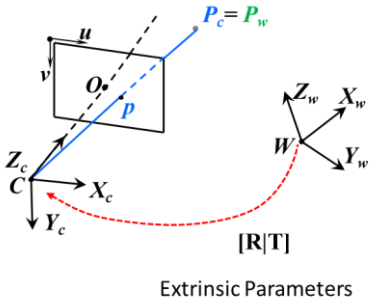

Stereo Rectification (2/5)

We can therefore write the Perspective Equation for the Left and Right cameras. For generality, we assume that Left and Right cameras have different intrinsic parameters (see also illustration below).

Left camera

$$\lambda_L \begin{bmatrix} u_L \\ v_L \\ 1 \end{bmatrix} = K_L R_L^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C_L \right)$$

Right camera

$$\lambda_R \begin{bmatrix} u_R \\ v_R \\ 1 \end{bmatrix} = K_R R_R^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C_R \right)$$

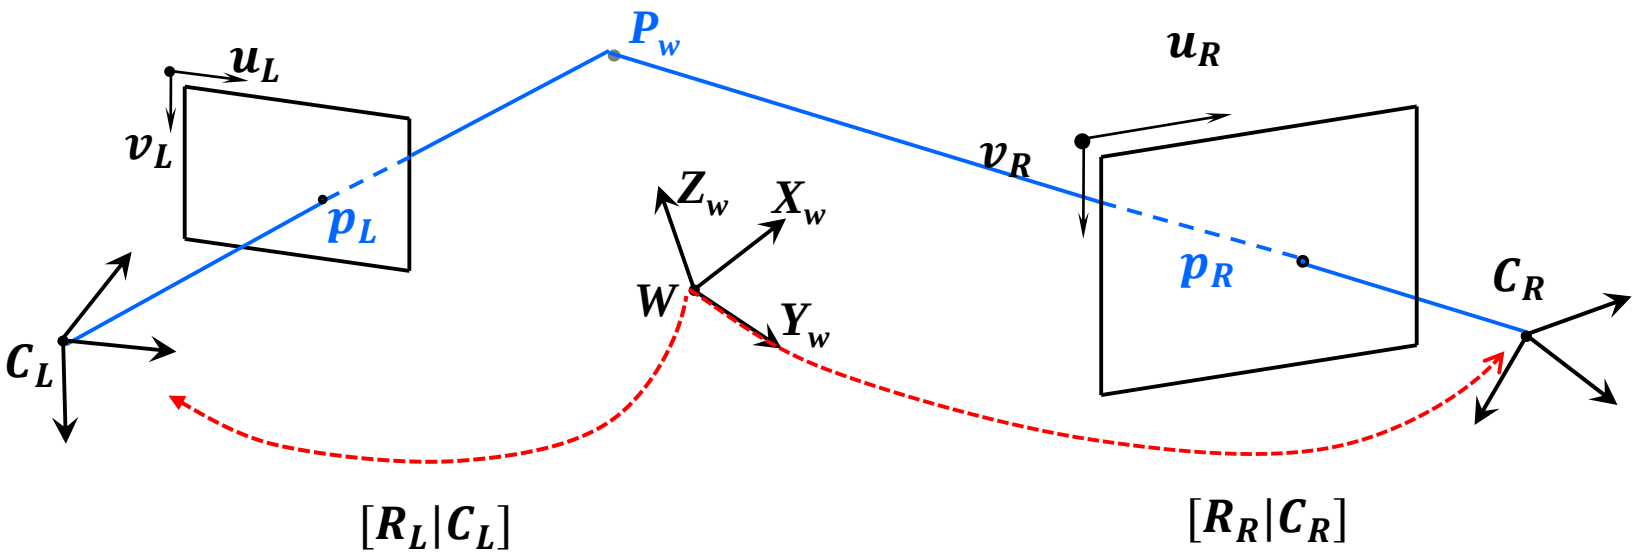

Stereo Rectification (3/5)

The goal of stereo rectification is to warp the left and right camera images such that their focal planes are coplanar and their intrinsic parameters are identical.

Old Left camera

$$\lambda_L \begin{bmatrix} u_L \\ v_L \\ 1 \end{bmatrix} = K_L R_L^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C_L \right)$$

New Left camera

$$\rightarrow \quad \bar{\lambda}_L \begin{bmatrix} \bar{u}_L \\ \bar{v}_L \\ 1 \end{bmatrix} = \bar{K} \bar{R}^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C_L \right)$$

Old Right camera

$$\lambda_R \begin{bmatrix} u_R \\ v_R \\ 1 \end{bmatrix} = K_R R_R^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C_R \right)$$

New Right camera

$$\rightarrow \quad \bar{\lambda}_R \begin{bmatrix} \bar{u}_R \\ \bar{v}_R \\ 1 \end{bmatrix} = \bar{K} \bar{R}^{-1} \left(\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix} - C_R \right)$$

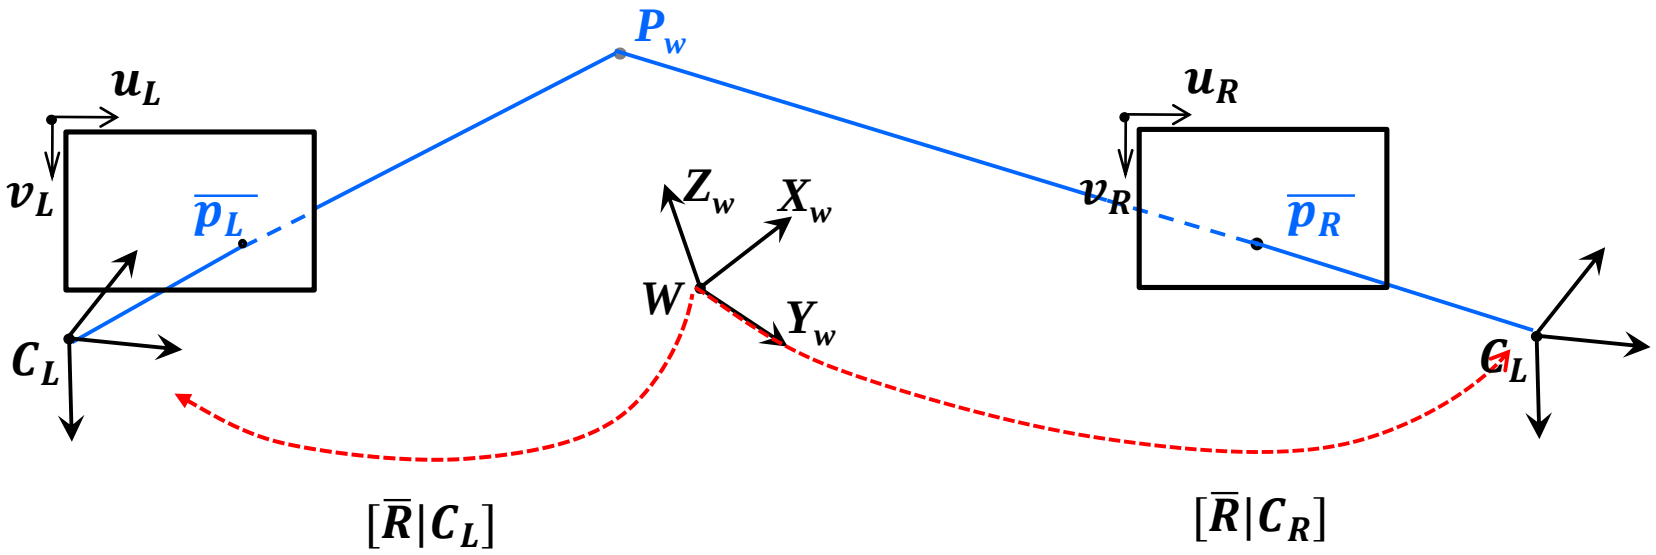

Stereo Rectification (4/5)

By solving for $\begin{bmatrix} X_w \\ Y_w \\ Z_w \end{bmatrix}$ for each camera, we can compute the Homography (or Warping) that needs to be applied to rectify each camera image

$$\bar{\lambda}_L \begin{bmatrix} \bar{u}_L \\ \bar{v}_L \\ 1 \end{bmatrix} = \lambda_L \underbrace{\bar{K}\,\bar{R}^{-1}_{\mathrm{L}}\,\mathrm{K}^{-1}_{\mathrm{L}}}_{\text{Homography of Left Camera}} \begin{bmatrix} u_L \\ v_L \\ 1 \end{bmatrix}$$

$$\bar{\lambda}_R \begin{bmatrix} \bar{u}_R \\ \bar{v}_R \\ 1 \end{bmatrix} = \lambda_R \underbrace{\bar{K}\,\bar{R}^{-1}_{\mathrm{R}}\,\mathrm{K}^{-1}_{\mathrm{R}}}_{\text{Homography of Right Camera}} \begin{bmatrix} u_R \\ v_R \\ 1 \end{bmatrix}$$

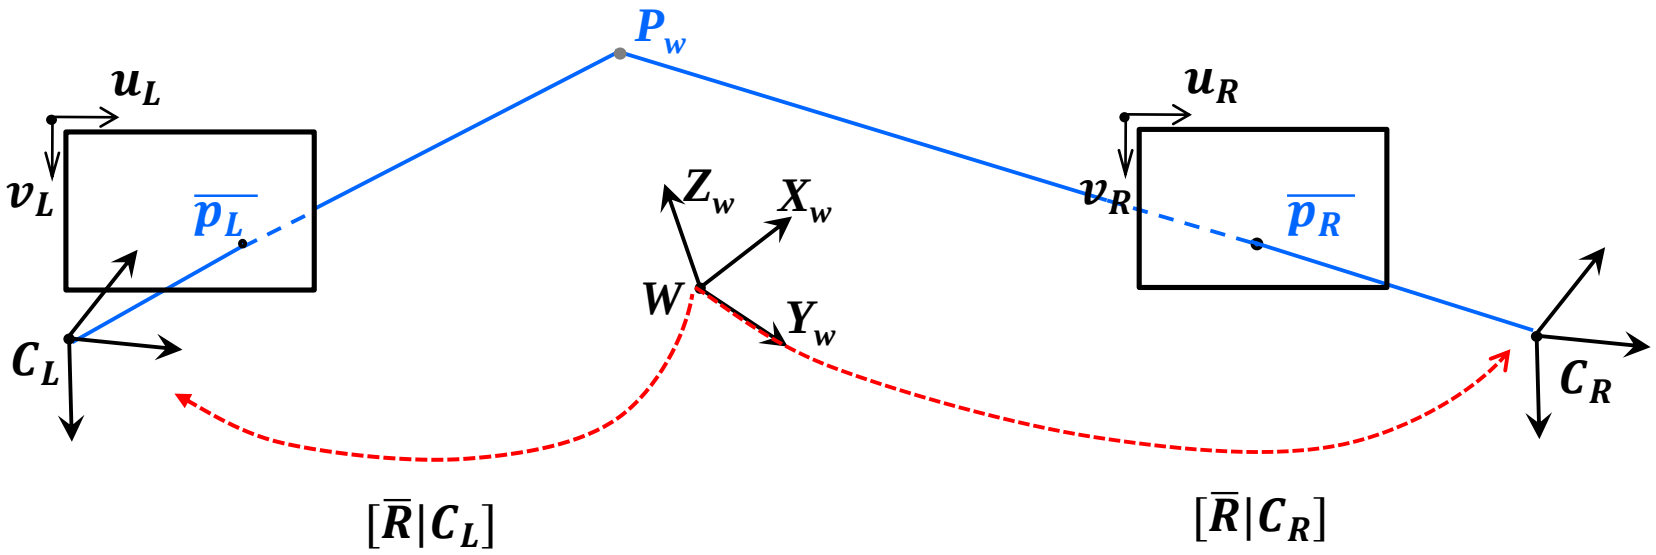

Stereo Rectification (5/5)

How do we chose the new $\bar{K}$ and $\bar{R}$? A good choice is to impose that:

$$\bar{K} = (K_L + K_R)/2$$

$$\bar{R} = [\bar{r_1}, \bar{r_2}, \bar{r_3}]$$

with $\bar{r_1}, \bar{r_2}, \bar{r_3}$ being the column vectors of $\bar{R}$, where

$$\bar{r_1} = \frac{C_2 - C_1}{\|C_2 - C_1\|}$$

$\bar{r_2} = r_3 \times \bar{r_1}$, where r_3 is the 3rd column of the rotation matrix of the left camera, i.e., R_L

$\bar{r_3} = \bar{r_1} \times \bar{r_2}$

More details can be found in the paper "A Compact Algorithm for Rectification of Stereo Pairs"

Stereo Rectification: example

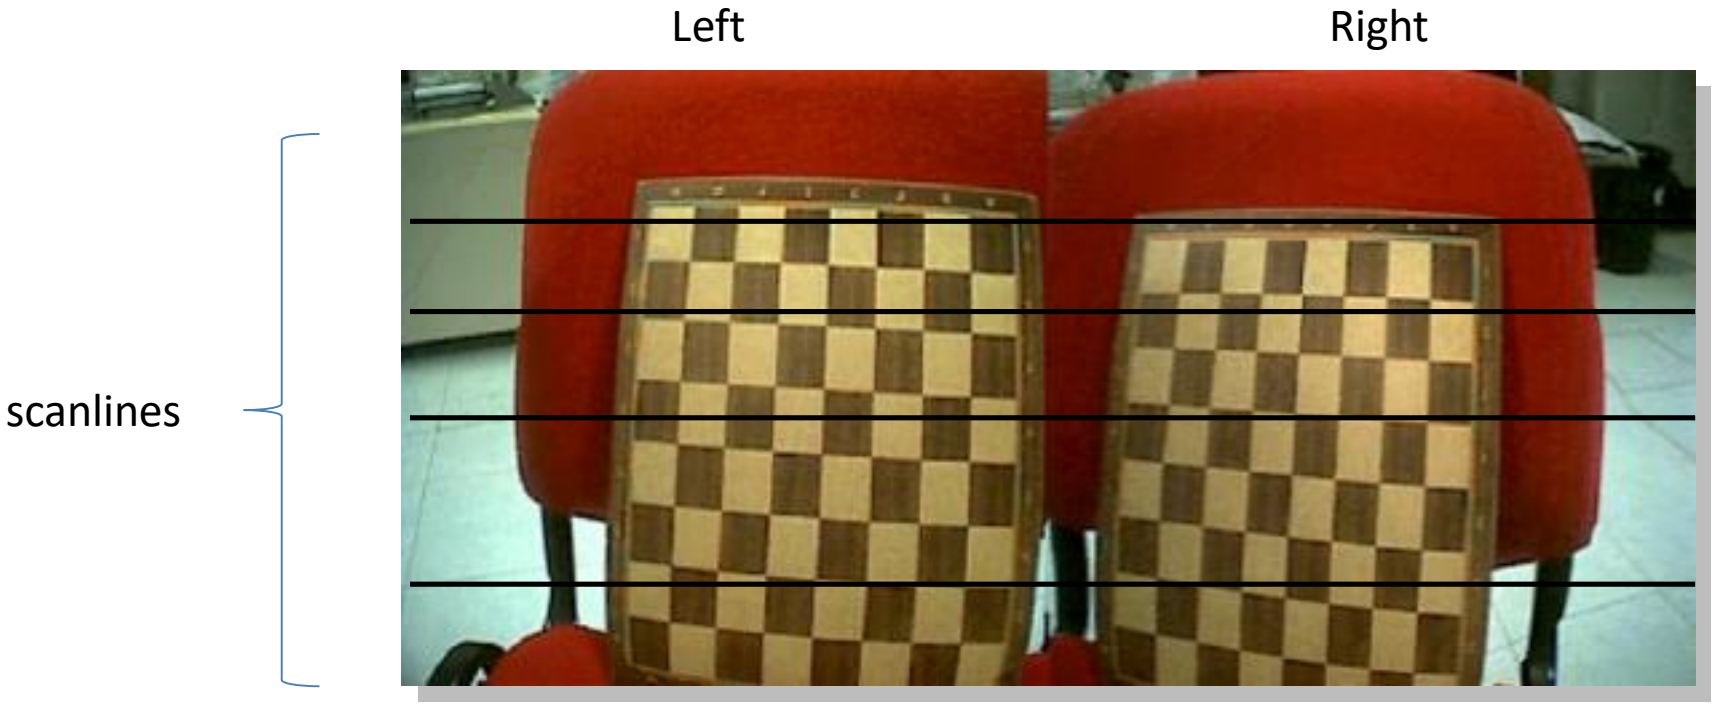

Stereo Rectification: example

- First, remove radial distortion (use bilinear interpolation (see lect. 06))

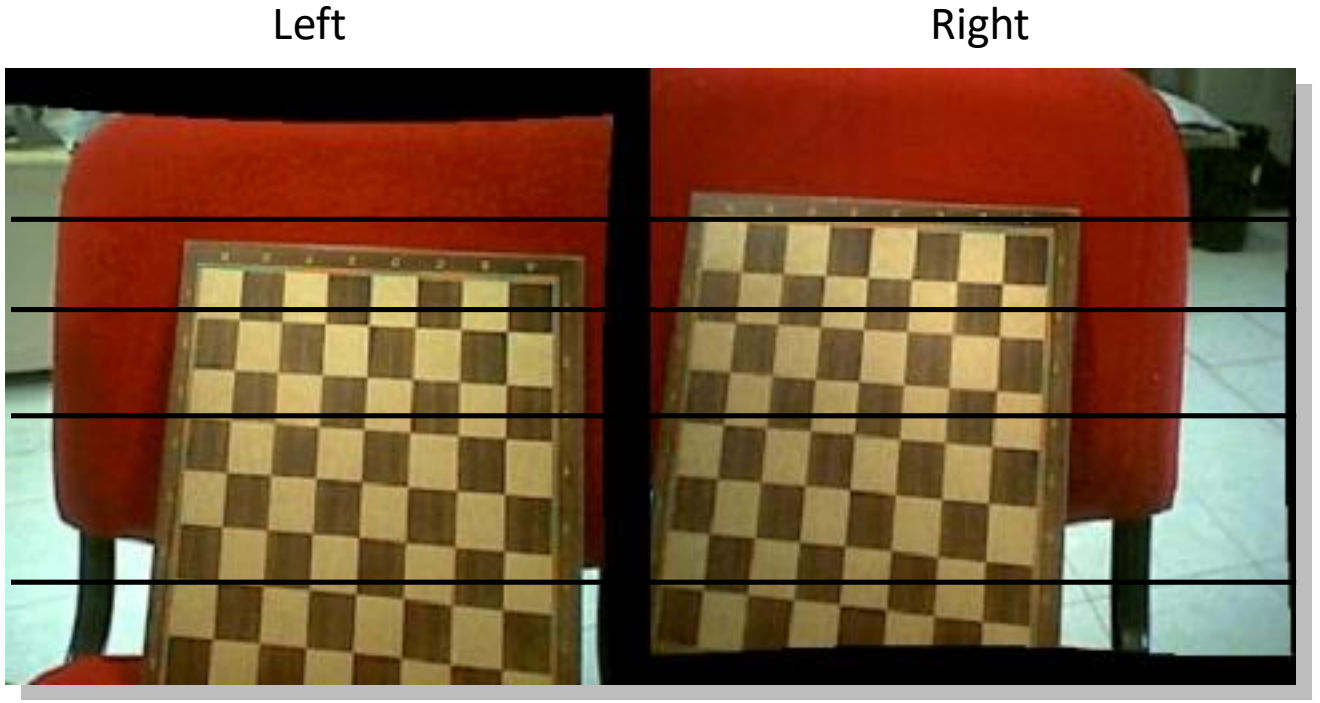

Stereo Rectification: example

- First, remove radial distortion (use bilinear interpolation (see lect. 06))
- Then, compute homographies and rectify (use bilinear interpolation)

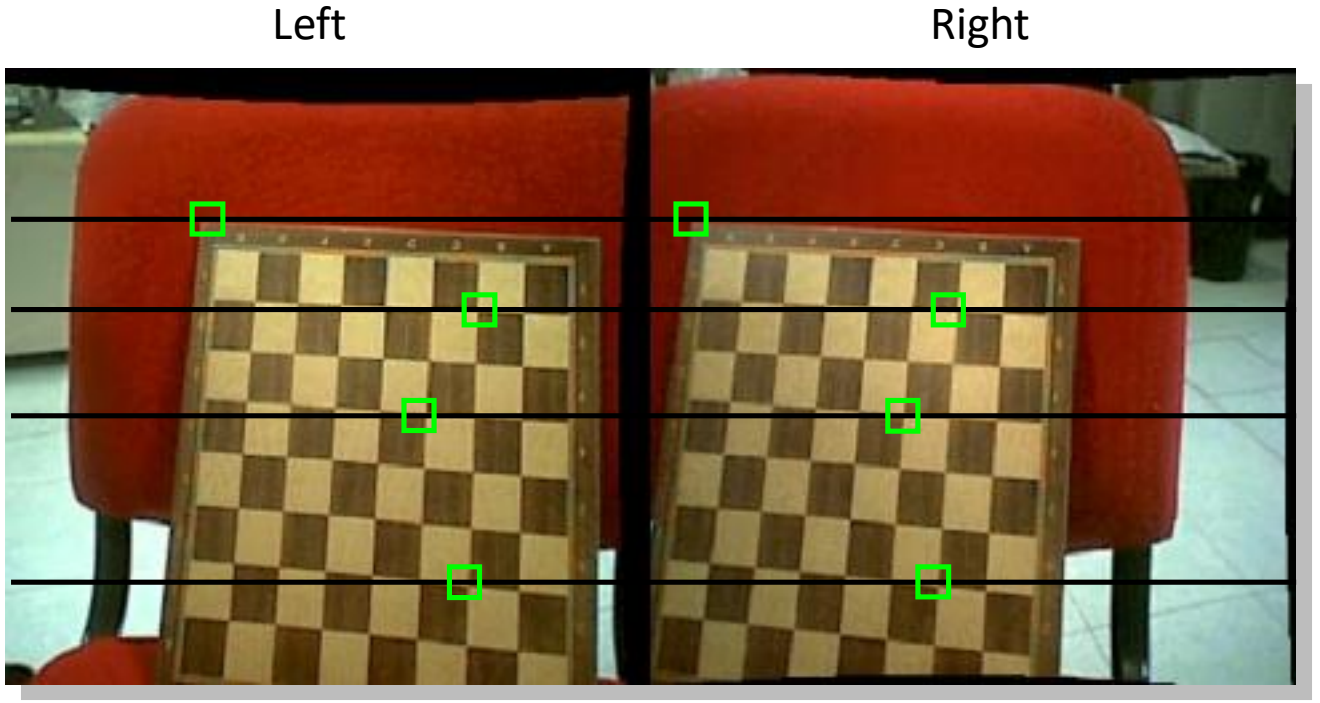

Stereo Rectification: example

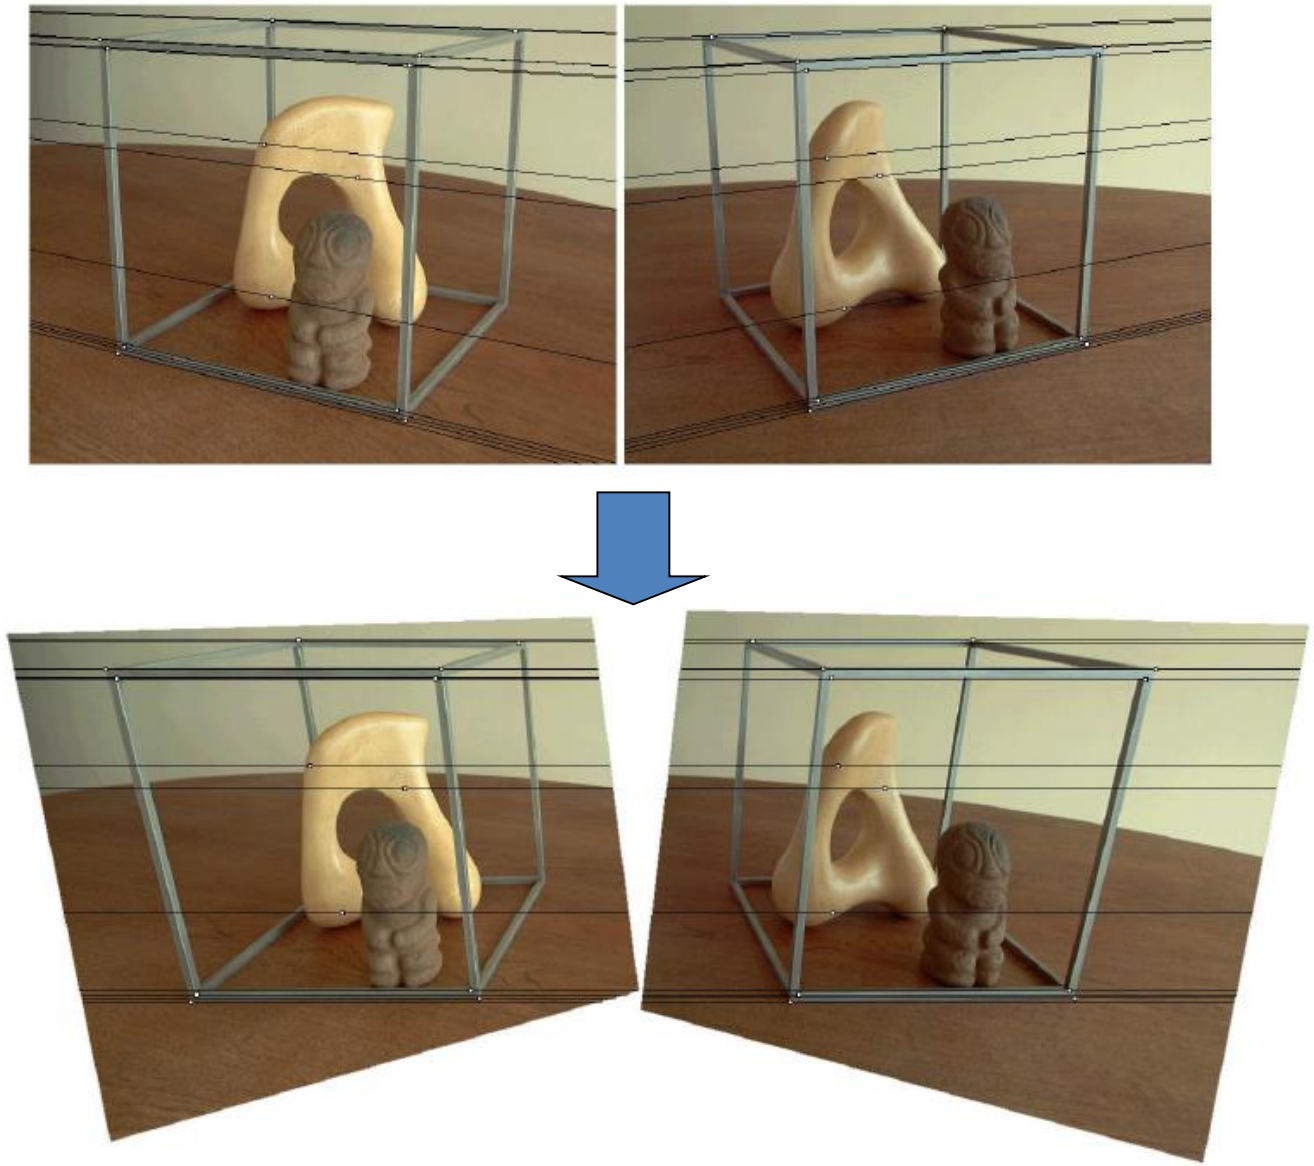

Stereo Vision

- Simplified case
- General case
- Correspondence problem (continued)
- Stereo rectification
- Triangulation

Correspondence problem

- Now that the left and right images are rectified, the correspondence search can be done along the same scanlines

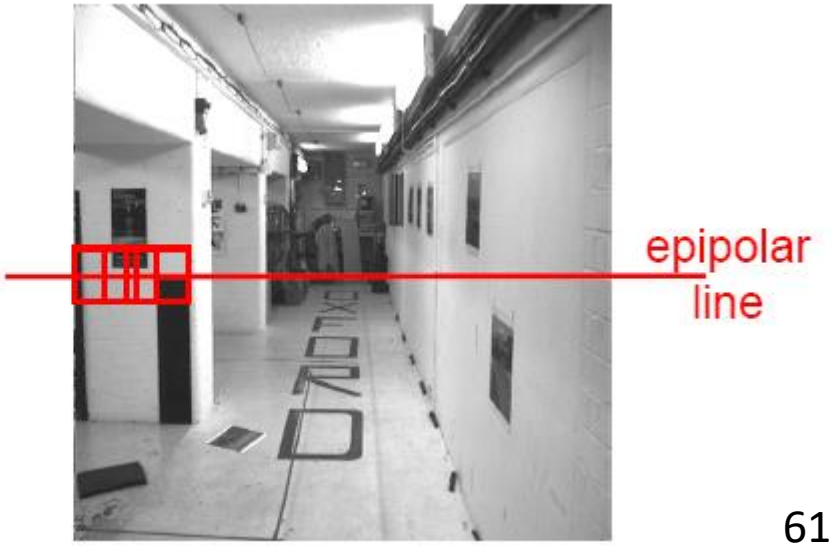

Correspondence problem

- To average noise effects, use a window around the point of interest
- Neighborhoods of corresponding points are similar in intensity patterns
- **Similarity measures**:
  - **(Z)NCC**
  - **(Z)SSD**
  - **(Z)SAD**
  - **Census Transform** (Census descriptor plus Hamming distance)

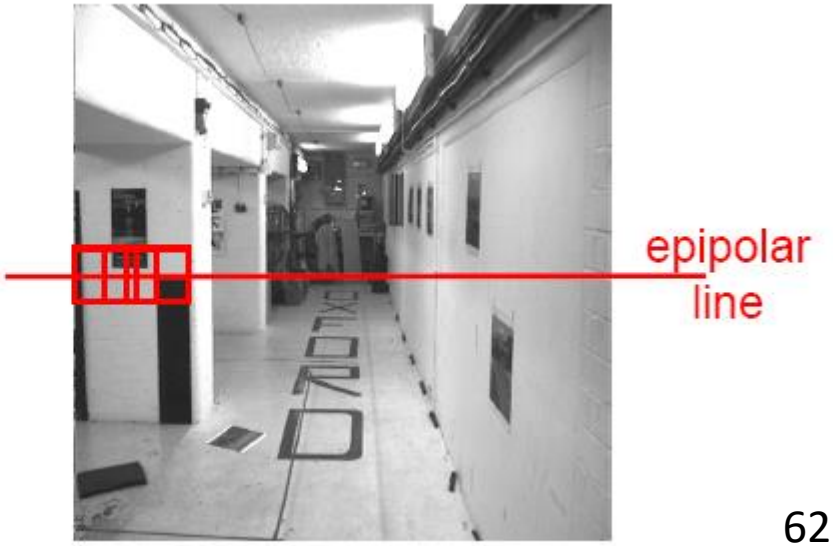

Correlation-based window matching

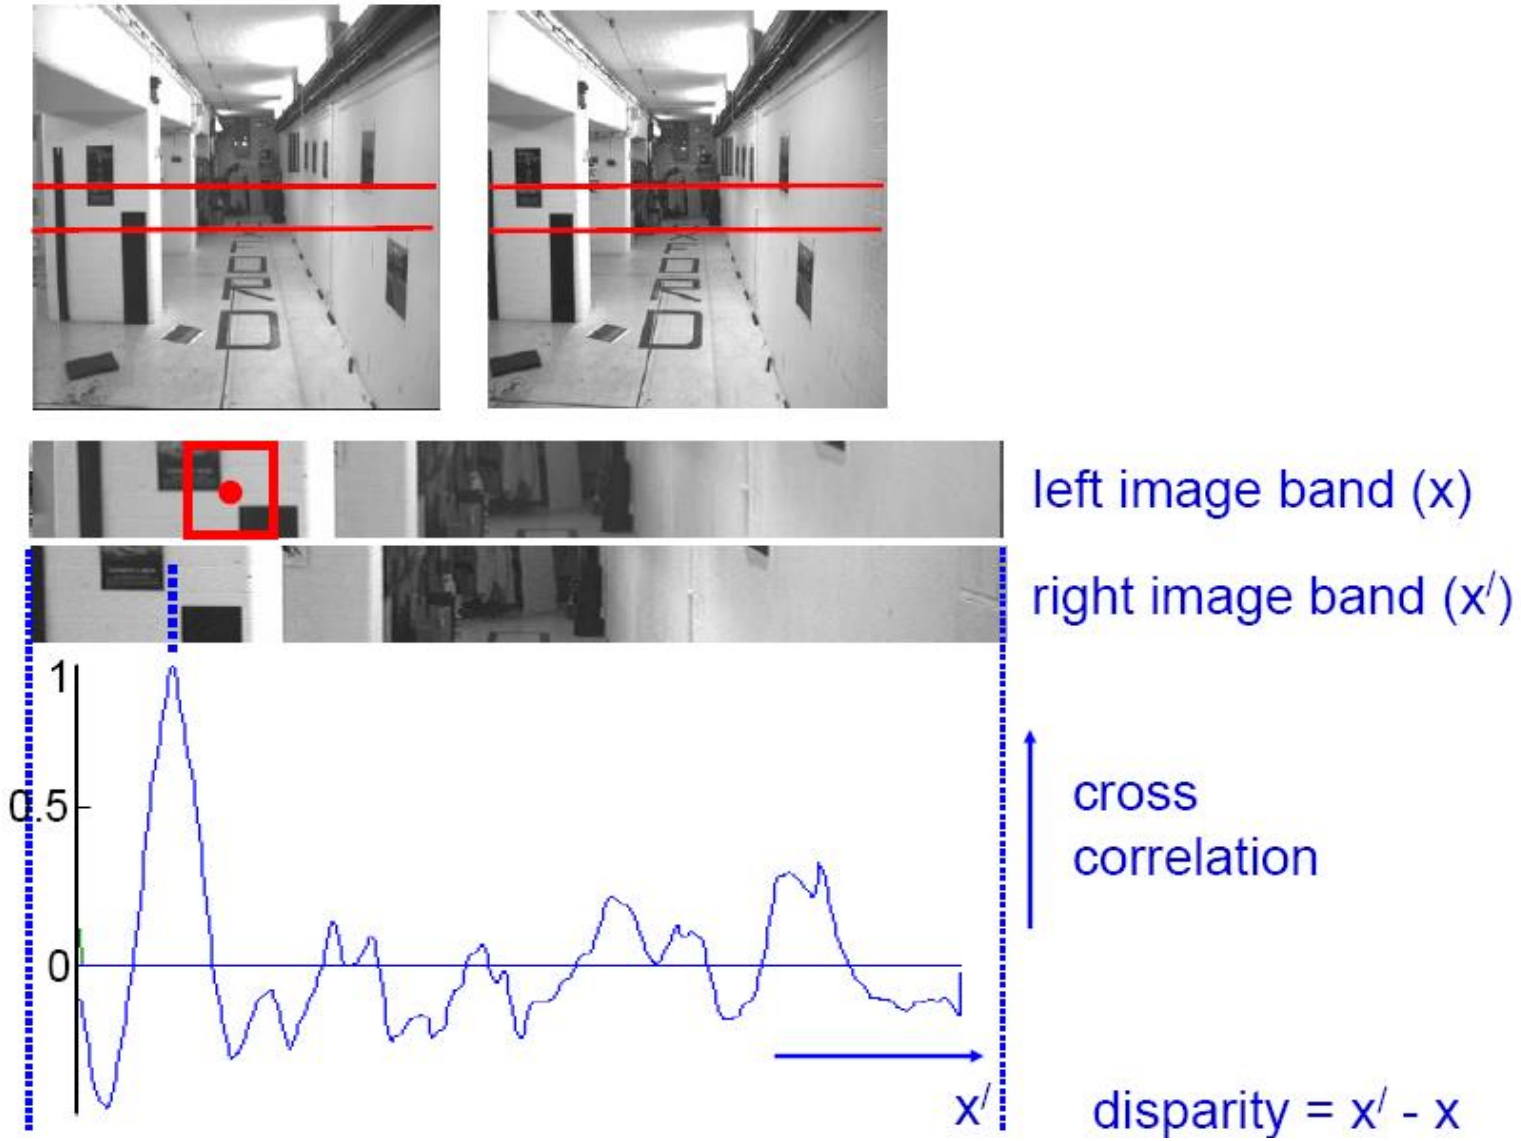

Correspondence Problems:
Textureless regions (**the aperture problem**)

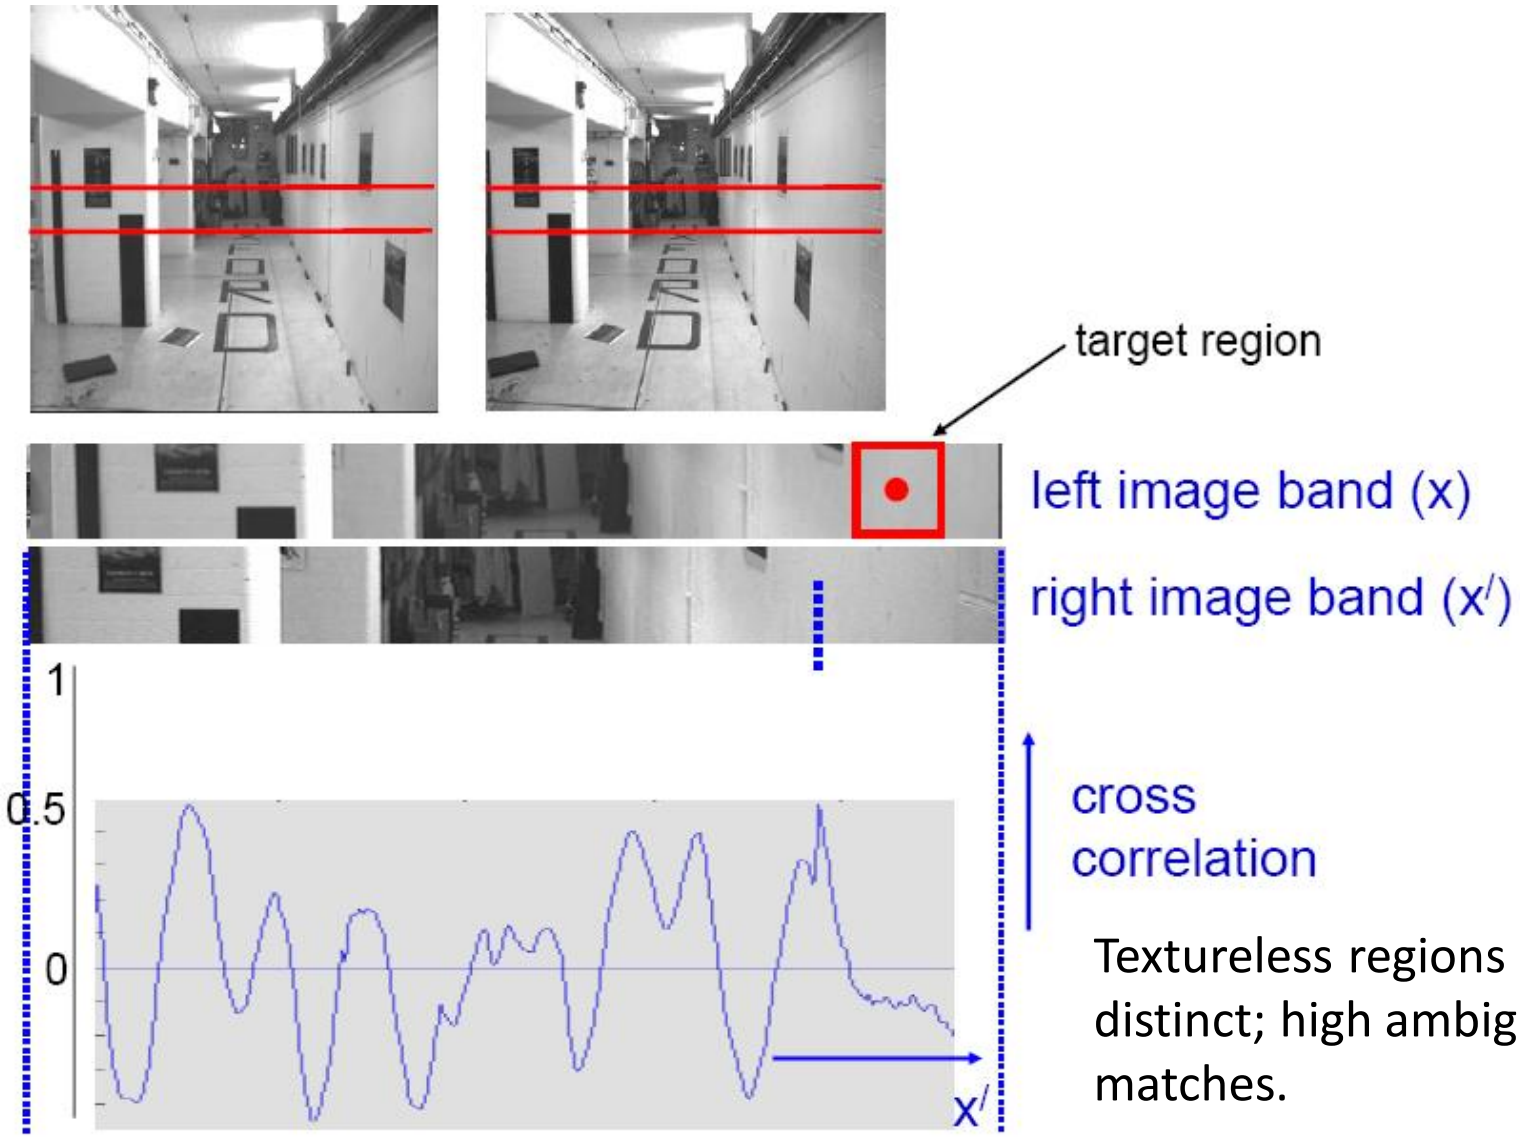

Textureless regions are non-distinct; high ambiguity for matches.

Solution: increase window size

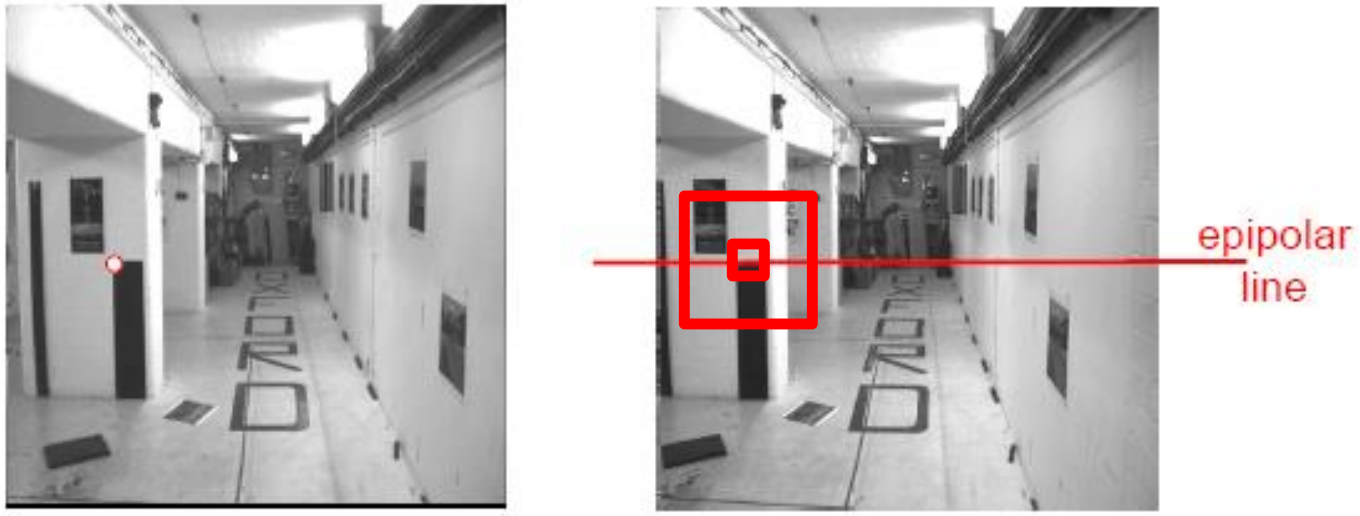

Effects of window size

W = 3

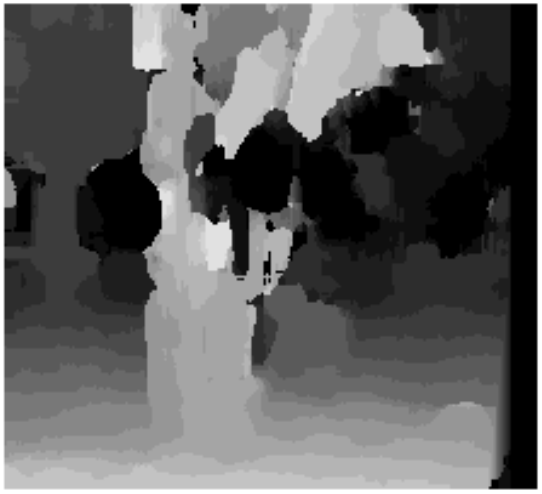

W = 20

- Smaller window
  - \+ More detail
  - – More noise
- Larger window
  - \+ Smoother disparity maps
  - – Less detail

Disparity map

Input to dense 3D reconstruction

1. For each pixel on the left image, find its corresponding point on the right image
2. Compute the disparity for each pair of correspondences
3. Visualize it in gray-scale or color coded image

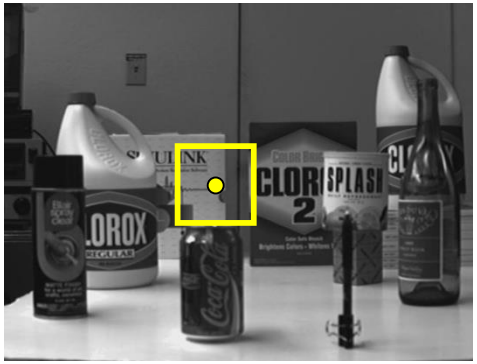

Left image

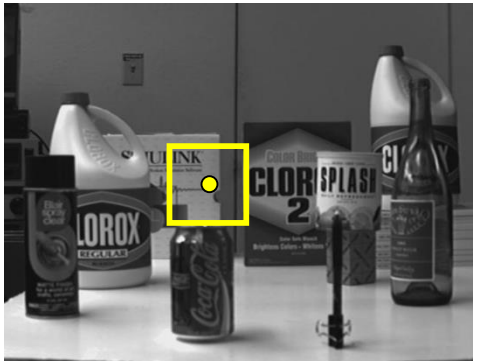

Right image

Close objects experience bigger disparity
⇨ appear brighter in disparity map

From Disparity Maps to Point Cloud

The depth Z can be computed from the disparity by recalling that $Z_P = \frac{bf}{u_l - u_r}$

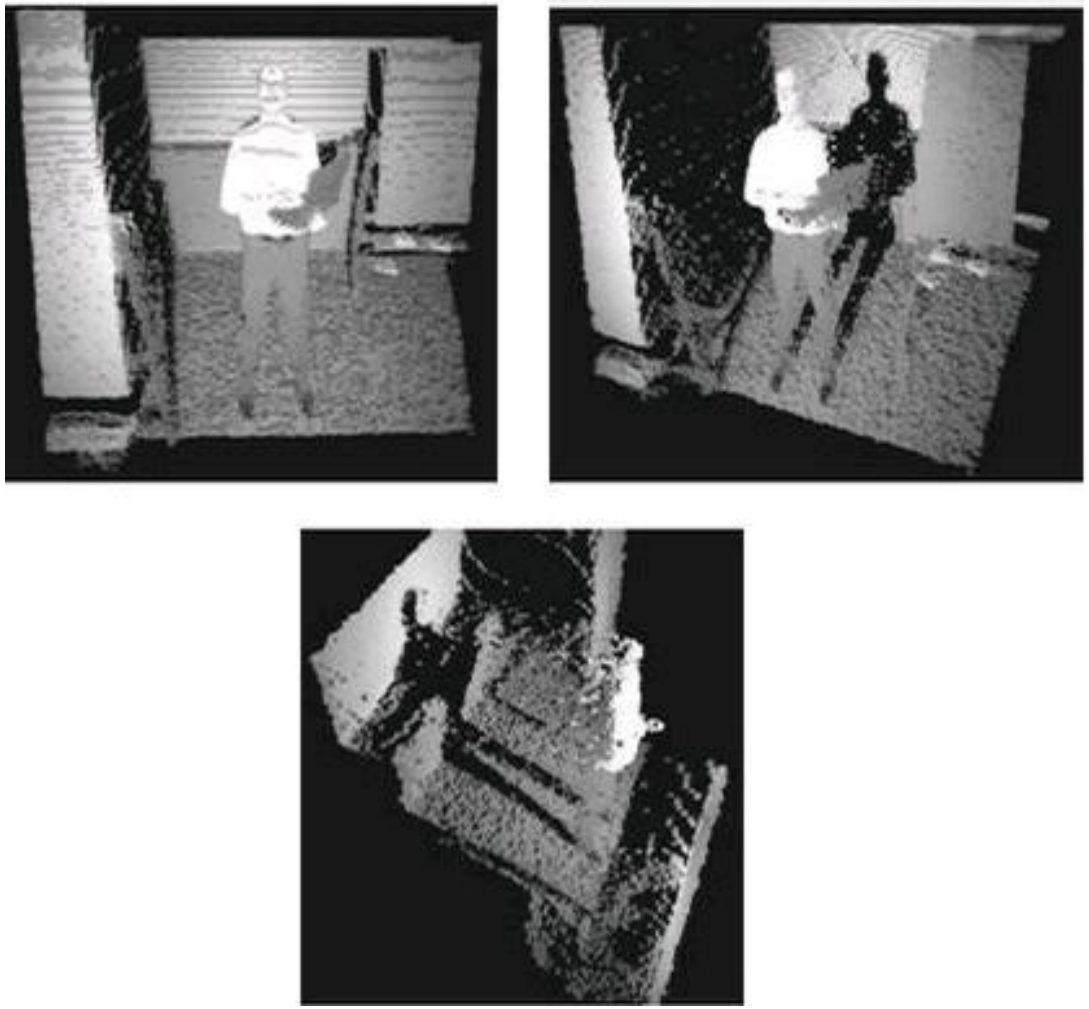

Accuracy

Data

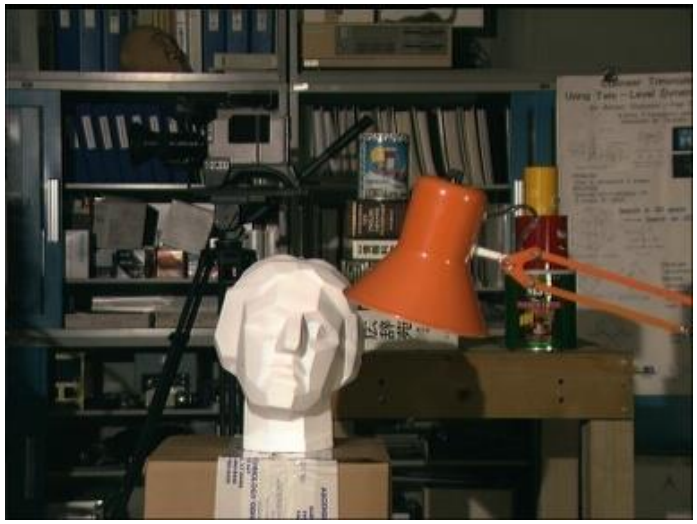

Window-based matching

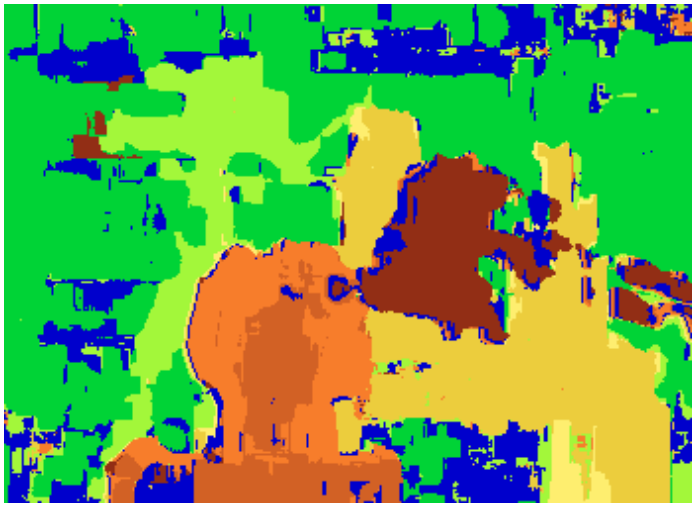

Ground truth

Challenges

Occlusions and repetitive patterns

Non-Lambertian surfaces (e.g., specularities), textureless surfaces

Correspondence Problems: Multiple matches

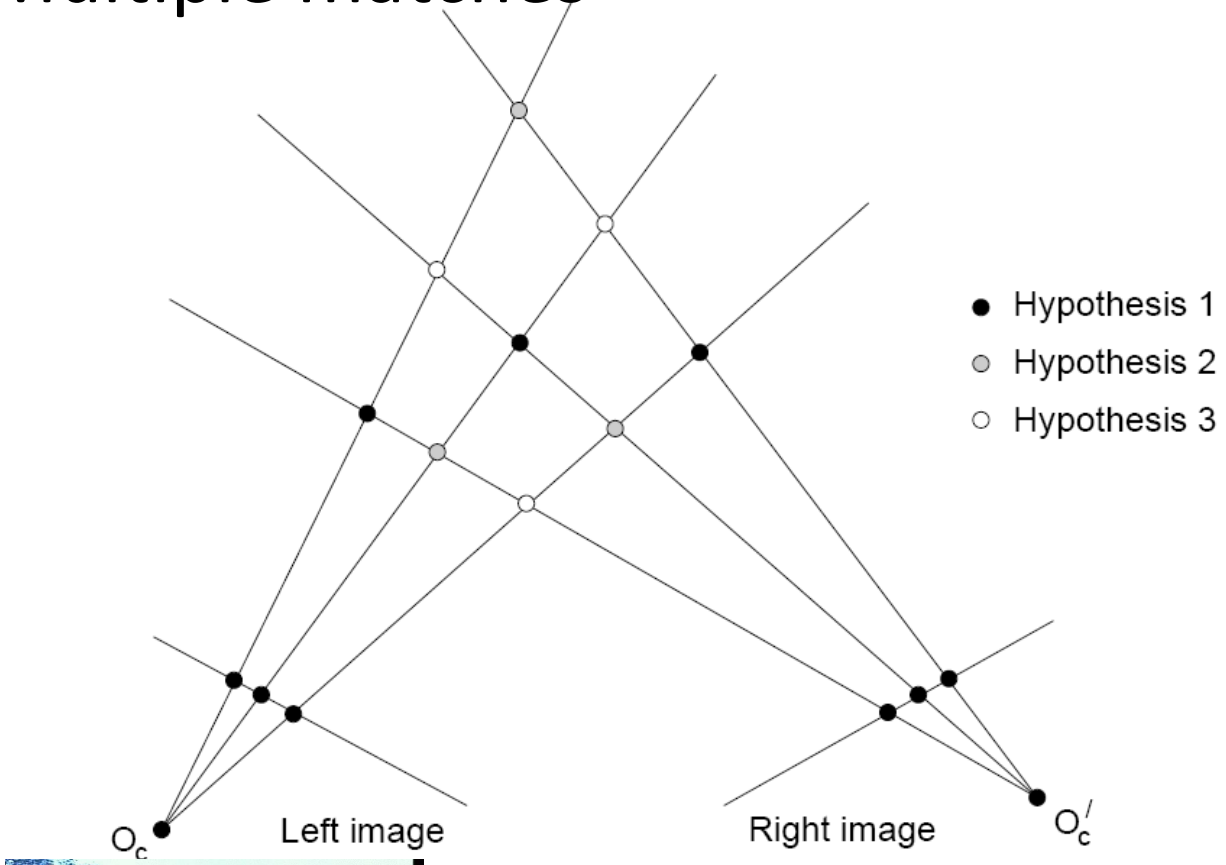

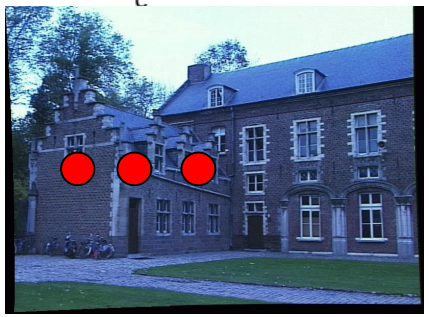

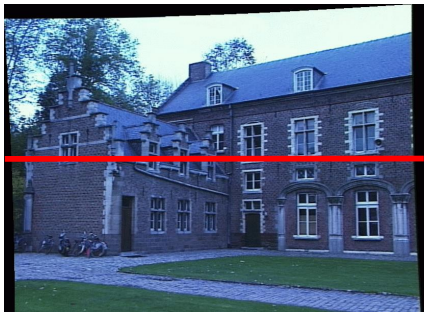

Multiple match hypotheses satisfy epipolar constraint, but which one is correct?

How can we improve window-based matching?

- Beyond the epipolar constraint, there are “soft” constraints to help identify corresponding points
  - Uniqueness
    - Only one match in right image for every point in left image
  - Ordering
    - Points on **same surface** will be in same order in both views
  - Disparity gradient
    - Disparity changes smoothly between points on the same surface

Better methods exist...

Using Deep Learning

Ground truth

Jia-Ren Chang Yong-Sheng Chen, Pyramid Stereo Matching Network, CVPR'18

For the latest and greatest:
http://vision.middlebury.edu/stereo/ and
http://www.cvlibs.net/datasets/kitti/eval_scene_flow.php?benchmark=stereo

Sparse Stereo Correspondence

- Restrict search to sparse set of detected features
- Feature matching
- Use epipolar geometry to narrow the search further

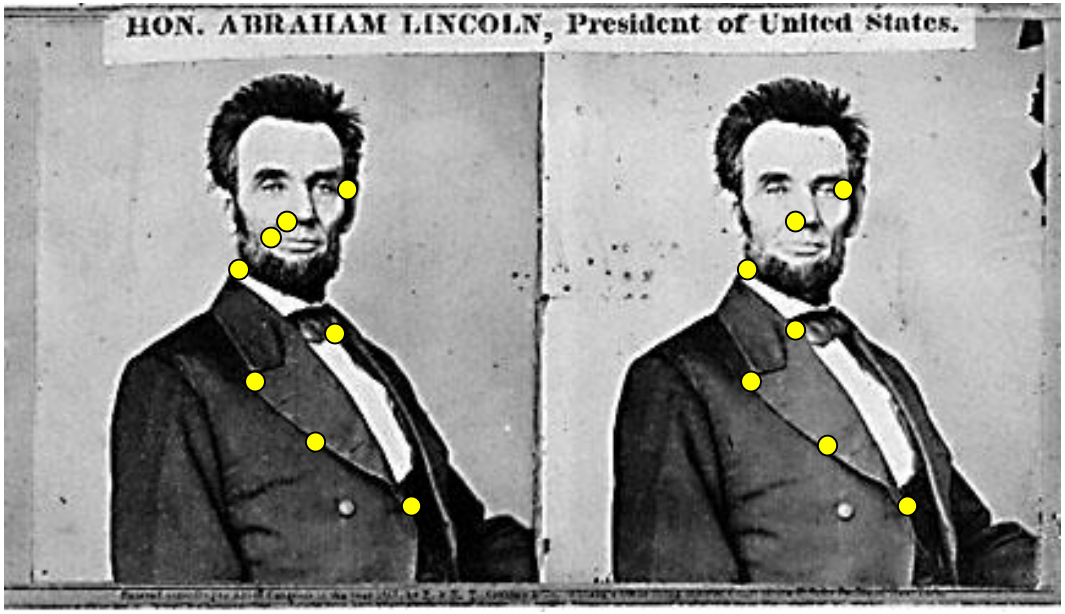

Things to Remember

- Disparity
- Triangulation: simplified and general case, linear and non linear approach
- Choosing the baseline
- Correspondence problem: epipoles, epipolar lines, epipolar plane
- Stereo rectification

- Readings:
  - Szeliski book: Chapter 11
  - Peter Corke book: Chapter 14.3
  - Autonomous Mobile Robot book: Chapter 4.2.5

Understanding Check

Are you able to answer the following questions?

- Can you relate Structure from Motion to 3D reconstruction? In what they differ?
- Can you define disparity in both the simplified and the general case?
- Can you provide a mathematical expression of depth as a function of the baseline, the disparity and the focal length?
- Can you apply error propagation to derive an expression for depth uncertainty? How can we improve the uncertainty?
- Can you analyze the effects of a large/small baseline?
- What is the closest depth that a stereo camera can measure?
- Are you able to show mathematically how to compute the intersection of two lines (linearly and non-linearly)?
- What is the geometric interpretation of the linear and non-linear approaches and what error do they minimize?
- Are you able to provide a definition of epipole, epipolar line and epipolar plane?
- Are you able to draw the epipolar lines for two converging cameras, for a forward motion situation, and for a side-moving camera?
- Are you able to define stereo rectification and to derive mathematically the rectifying homographies?
- How is the disparity map computed?
- How can one establish stereo correspondences with subpixel accuracy?
- Describe one or more simple ways to reject outliers in stereo correspondences.
- Is stereo vision the only way of estimating depth information? If not, are you able to list alternative options?